

Human activity recognition

This project is to build a model that predicts the human activities such as Walking, Walking_Upstairs, Walking_Downstairs, Sitting, Standing or Laying.

This dataset is collected from 30 persons(referred as subjects in this dataset), performing different activities with a smartphone to their waists. The data is recorded with the help of sensors (accelerometer and Gyroscope) in that smartphone. This experiment was video recorded to label the data manually.

How data was recorded

By using the sensors(Gyroscope and accelerometer) in a smartphone, they have captured '3-axial linear acceleration'(*tAcc-XYZ*) from accelerometer and '3-axial angular velocity' (*tGyro-XYZ*) from Gyroscope with several variations.

prefix 't' in those metrics denotes time.

suffix 'XYZ' represents 3-axial signals in X , Y, and Z directions.

Feature names

1. These sensor signals are preprocessed by applying noise filters and then sampled in fixed-width windows(sliding windows) of 2.56 seconds each with 50% overlap. ie., each window has 128 readings.
2. From Each window, a feature vector was obtained by calculating variables from the time and frequency domain.

In our dataset, each datapoint represents a window with different readings

3. The accelertion signal was saperated into Body and Gravity acceleration signals(***tBodyAcc-XYZ*** and ***tGravityAcc-XYZ***) using some low pass filter with corner frequency of 0.3Hz.
4. After that, the body linear acceleration and angular velocity were derived in time to obtian *jerk signals* (***tBodyAccJerk-XYZ*** and ***tBodyGyroJerk-XYZ***).
5. The magnitude of these 3-dimensional signals were calculated using the Euclidian norm. This magnitudes are represented as features with names like *tBodyAccMag*, *tGravityAccMag*, *tBodyAccJerkMag*, *tBodyGyroMag* and *tBodyGyroJerkMag*.
6. Finally, We've got frequency domain signals from some of the available signals by applying a FFT (Fast Fourier Transform). These signals obtained were labeled with **prefix 'f'** just like original signals with **prefix 't'**. These signals are labeled as ***fBodyAcc-XYZ***, ***fBodyGyroMag*** etc.,.
7. These are the signals that we got so far.

- *tBodyAcc-XYZ*
- *tGravityAcc-XYZ*
- *tBodyAccJerk-XYZ*
- *tBodyGyro-XYZ*
- *tBodyGyroJerk-XYZ*
- *tBodyAccMag*
- *tGravityAccMag*
- *tBodyAccJerkMag*
- *tBodyGyroMag*
- *tBodyGyroJerkMag*
- *fBodyAcc-XYZ*
- *fBodyAccJerk-XYZ*

- fBodyGyro-XYZ
- fBodyAccMag
- fBodyAccJerkMag
- fBodyGyroMag
- fBodyGyroJerkMag

8. We can estimate some set of variables from the above signals. ie., We will estimate the following properties on each and every signal that we recorded so far.

- **mean()**: Mean value
- **std()**: Standard deviation
- **mad()**: Median absolute deviation
- **max()**: Largest value in array
- **min()**: Smallest value in array
- **sma()**: Signal magnitude area
- **energy()**: Energy measure. Sum of the squares divided by the number of values.
- **iqr()**: Interquartile range
- **entropy()**: Signal entropy
- **arCoeff()**: Autoregression coefficients with Burg order equal to 4
- **correlation()**: correlation coefficient between two signals
- **maxInds()**: index of the frequency component with largest magnitude
- **meanFreq()**: Weighted average of the frequency components to obtain a mean frequency
- **skewness()**: skewness of the frequency domain signal
- **kurtosis()**: kurtosis of the frequency domain signal
- **bandsEnergy()**: Energy of a frequency interval within the 64 bins of the FFT of each window.
- **angle()**: Angle between two vectors.

9. We can obtain some other vectors by taking the average of signals in a single window sample. These are used on the angle() variable'

- gravityMean
- tBodyAccMean
- tBodyAccJerkMean
- tBodyGyroMean
- tBodyGyroJerkMean

Y_Labels(Encoded)

- In the dataset, Y_labels are represented as numbers from 1 to 6 as their identifiers.
 - WALKING as **1**
 - WALKING_UPSTAIRS as **2**
 - WALKING_DOWNSTAIRS as **3**
 - SITTING as **4**
 - STANDING as **5**
 - LAYING as **6**

Train and test data were separated

- The readings from **70%** of the volunteers were taken as **training data** and remaining **30%** subjects recordings were taken for **test data**

Data

- All the data is present in 'UCI_HAR_dataset/' folder in present working directory.
 - Feature names are present in 'UCI_HAR_dataset/features.txt'
 - **Train Data**
 - 'UCI_HAR_dataset/train/X_train.txt'
 - 'UCI_HAR_dataset/train/subject_train.txt'
 - 'UCI_HAR_dataset/train/y_train.txt'
 - **Test Data**
 - 'UCI_HAR_dataset/test/X_test.txt'
 - 'UCI_HAR_dataset/test/subject_test.txt'
 - 'UCI_HAR_dataset/test/y_test.txt'

Data Size :

27 MB

Quick overview of the dataset :

- Accelerometer and Gyroscope readings are taken from 30 volunteers(referred as subjects) while performing the following 6 Activities.
 1. Walking
 2. WalkingUpstairs
 3. WalkingDownstairs
 4. Standing
 5. Sitting
 6. Lying.
- Readings are divided into a window of 2.56 seconds with 50% overlapping.
- Accelerometer readings are divided into gravity acceleration and body acceleration readings, which has x,y and z components each.
- Gyroscope readings are the measure of angular velocities which has x,y and z components.
- Jerk signals are calculated for BodyAcceleration readings.
- Fourier Transforms are made on the above time readings to obtain frequency readings.
- Now, on all the base signal readings., mean, max, mad, sma, arcoefficient, engerybands,entropy etc., are calculated for each window.
- We get a feature vector of 561 features and these features are given in the dataset.
- Each window of readings is a datapoint of 561 features.

Problem Framework

- 30 subjects(volunteers) data is randomly split to 70%(21) test and 30%(7) train data.
- Each datapoint corresponds one of the 6 Activities.

Problem Statement

- Given a new datapoint we have to predict the Activity

In [1]:

```
import numpy as np
import pandas as pd

# get the features from the file features.txt
features = list()
with open('UCI_HAR_Dataset/features.txt') as f:
    features = [line.split()[1] for line in f.readlines()]
print('No of Features: {}'.format(len(features)))
```

No of Features: 561

Obtain the train data

In [2]:

```
# get the data from txt files to pandas dataframe
X_train = pd.read_csv('UCI_HAR_dataset/train/X_train.txt', delim_whitespace=True, header=None, names=features)

# add subject column to the dataframe
X_train['subject'] = pd.read_csv('UCI_HAR_dataset/train/subject_train.txt', header=None, squeeze=True)

y_train = pd.read_csv('UCI_HAR_dataset/train/y_train.txt', names=['Activity'], squeeze=True)
y_train_labels = y_train.map({1: 'WALKING', 2: 'WALKING_UPSTAIRS', 3: 'WALKING_DOWNSTAIRS',
                               4: 'SITTING', 5: 'STANDING', 6: 'LAYING'})


# put all columns in a single dataframe
train = X_train
train['Activity'] = y_train
train['ActivityName'] = y_train_labels
train.sample()
```

D:\installed\Anaconda3\lib\site-packages\pandas\io\parsers.py:678: UserWarning: Duplicate names specified. This will raise an error in the future.
 return _read(filepath_or_buffer, kwds)

Out[2]:

	tBodyAcc-mean()-X	tBodyAcc-mean()-Y	tBodyAcc-mean()-Z	tBodyAcc-std()-X	tBodyAcc-std()-Y	tBodyAcc-std()-Z	tBodyAccmad()
6015	0.2797	-0.004397	-0.10952	0.359081	0.119909	-0.177541	0.337963

1 rows × 564 columns



In [3]:

train.shape

Out[3]:

(7352, 564)

Obtain the test data

In [4]:

```
# get the data from txt files to pandas dataframe
X_test = pd.read_csv('UCI_HAR_dataset/test/X_test.txt', delim_whitespace=True, header=None, names=features)

# add subject column to the dataframe
X_test['subject'] = pd.read_csv('UCI_HAR_dataset/test/subject_test.txt', header=None, squeeze=True)

# get y labels from the txt file
y_test = pd.read_csv('UCI_HAR_dataset/test/y_test.txt', names=['Activity'], squeeze=True)
y_test_labels = y_test.map({1: 'WALKING', 2: 'WALKING_UPSTAIRS', 3: 'WALKING_DOWNSTAIRS', 4: 'SITTING', 5: 'STANDING', 6: 'LAYING'})

# put all columns in a single dataframe
test = X_test
test['Activity'] = y_test
test['ActivityName'] = y_test_labels
test.sample()
```


D:\installed\Anaconda3\lib\site-packages\pandas\io\parsers.py:678: UserWarning: Duplicate names specified. This will raise an error in the future.

```
return _read(filepath_or_buffer, kwds)
```

Out[4]:

	tBodyAcc-mean()-X	tBodyAcc-mean()-Y	tBodyAcc-mean()-Z	tBodyAcc-std()-X	tBodyAcc-std()-Y	tBodyAcc-std()-Z	tBodyAccmad()
2261	0.279196	-0.018261	-0.103376	-0.996955	-0.982959	-0.988239	-0.9972

1 rows × 564 columns



In [5]:

test.shape

Out[5]:

(2947, 564)

Data Cleaning

1. Check for Duplicates

In [6]:

```
print('No of duplicates in train: {}'.format(sum(train.duplicated())))  
print('No of duplicates in test : {}'.format(sum(test.duplicated())))
```

```
No of duplicates in train: 0  
No of duplicates in test : 0
```

2. Checking for NaN/null values

In [7]:

```
print('We have {} NaN/Null values in train'.format(train.isnull().values.sum()))  
print('We have {} NaN/Null values in test'.format(test.isnull().values.sum()))
```

```
We have 0 NaN/Null values in train  
We have 0 NaN/Null values in test
```


3. Check for data imbalance

In [8]:

```
import matplotlib.pyplot as plt  
import seaborn as sns  
  
sns.set_style('whitegrid')  
plt.rcParams['font.family'] = 'Dejavu Sans'
```

In [9]:


```
plt.figure(figsize=(16,8))  
plt.title('Data provided by each user', fontsize=20)  
sns.countplot(x='subject',hue='ActivityName', data = train)  
plt.show()
```



We have got almost same number of reading from all the subjects

In [10]:

```
plt.title('No of Datapoints per Activity', fontsize=15)
sns.countplot(train.ActivityName)
plt.xticks(rotation=90)
plt.show()
```



Observation

Our data is well balanced (almost)

4. Changing feature names

In [11]:

```
columns = train.columns

# Removing '()' from column names
columns = columns.str.replace('[(\)]', '')
columns = columns.str.replace('[-]', '')
columns = columns.str.replace('[,]', '')

train.columns = columns
test.columns = columns

test.columns
```

Out[11]:

```
Index(['tBodyAccmeanX', 'tBodyAccmeanY', 'tBodyAccmeanZ', 'tBodyAccstdX',
      'tBodyAccstdY', 'tBodyAccstdZ', 'tBodyAccmadX', 'tBodyAccmadY',
      'tBodyAccmadZ', 'tBodyAccmaxX',
      ...,
      'angletBodyAccMeangravity', 'angletBodyAccJerkMeangravityMean',
      'angletBodyGyroMeangravityMean', 'angletBodyGyroJerkMeangravityMea
n',
      'angleXgravityMean', 'angleYgravityMean', 'angleZgravityMean',
      'subject', 'Activity', 'ActivityName'],
      dtype='object', length=564)
```

5. Save this dataframe in a csv files

In [13]:

```
train.to_csv('UCI_HAR_Dataset/csv_files/train.csv', index=False)
test.to_csv('UCI_HAR_Dataset/csv_files/test.csv', index=False)
```

Exploratory Data Analysis

"Without domain knowledge EDA has no meaning, without EDA a problem has no soul."

1. Featuring Engineering from Domain Knowledge

- **Static and Dynamic Activities**

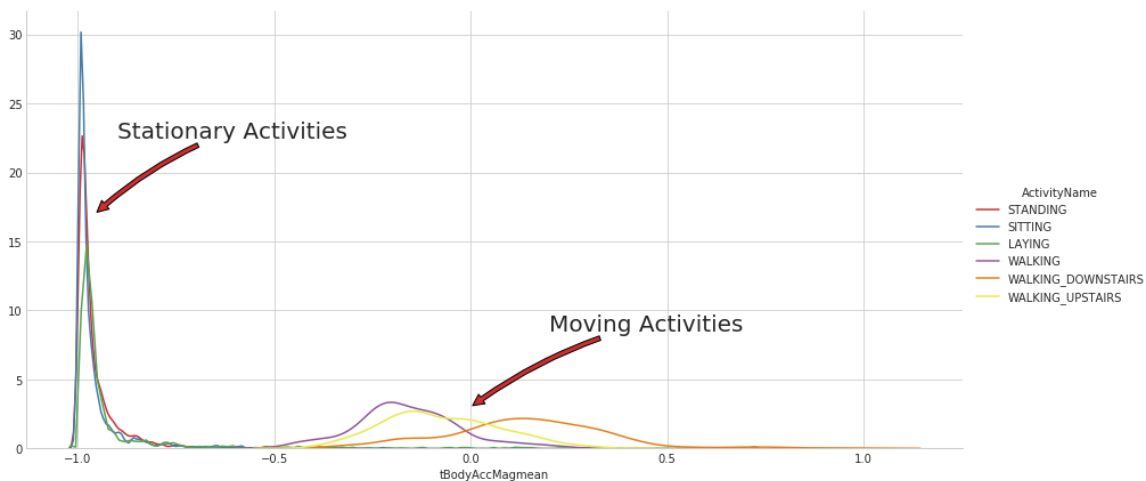
- In static activities (sit, stand, lie down) motion information will not be very useful.
- In the dynamic activities (Walking, WalkingUpstairs, WalkingDownstairs) motion info will be significant.

2. Stationary and Moving activities are completely different

In [14]:

```
sns.set_palette("Set1", desat=0.80)
facetgrid = sns.FacetGrid(train, hue='ActivityName', size=6, aspect=2)
facetgrid.map(sns.distplot, 'tBodyAccMagmean', hist=False)\
    .add_legend()
plt.annotate("Stationary Activities", xy=(-0.956,17), xytext=(-0.9, 23), size=20,\
    va='center', ha='left',\
    arrowprops=dict(arrowstyle="simple",connectionstyle="arc3,rad=0.1"))

plt.annotate("Moving Activities", xy=(0,3), xytext=(0.2, 9), size=20,\
    va='center', ha='left',\
    arrowprops=dict(arrowstyle="simple",connectionstyle="arc3,rad=0.1"))
plt.show()
```




In [15]:

```
# for plotting purposes taking datapoints of each activity to a different dataframe
df1 = train[train['Activity']==1]
df2 = train[train['Activity']==2]
df3 = train[train['Activity']==3]
df4 = train[train['Activity']==4]
df5 = train[train['Activity']==5]
df6 = train[train['Activity']==6]

plt.figure(figsize=(14,7))
plt.subplot(2,2,1)
plt.title('Stationary Activities(Zoomed in)')
sns.distplot(df4['tBodyAccMagmean'],color = 'r',hist = False, label = 'Sitting')
sns.distplot(df5['tBodyAccMagmean'],color = 'm',hist = False,label = 'Standing')
sns.distplot(df6['tBodyAccMagmean'],color = 'c',hist = False, label = 'Laying')
plt.axis([-1.01, -0.5, 0, 35])
plt.legend(loc='center')

plt.subplot(2,2,2)
plt.title('Moving Activities')
sns.distplot(df1['tBodyAccMagmean'],color = 'red',hist = False, label = 'Walking')
sns.distplot(df2['tBodyAccMagmean'],color = 'blue',hist = False,label = 'Walking Up')
sns.distplot(df3['tBodyAccMagmean'],color = 'green',hist = False, label = 'Walking down')
plt.legend(loc='center right')


plt.tight_layout()
plt.show()
```



3. Magnitude of an acceleration can saperate it well

In [16]:

```
plt.figure(figsize=(7,7))
sns.boxplot(x='ActivityName', y='tBodyAccMagmean', data=train, showfliers=False, saturation=1)
plt.ylabel('Acceleration Magnitude mean')
plt.axhline(y=-0.7, xmin=0.1, xmax=0.9, dashes=(5,5), c='g')
plt.axhline(y=-0.05, xmin=0.4, xmax=0.9, dashes=(5,5), c='m')
plt.xticks(rotation=90)
plt.show()
```




Observations:

- If tAccMean is < -0.8 then the Activities are either Standing or Sitting or Laying.
- If tAccMean is > -0.6 then the Activities are either Walking or WalkingDownstairs or WalkingUpstairs.
- If tAccMean > 0.0 then the Activity is WalkingDownstairs.
- We can classify 75% the Activity labels with some errors.

4. Position of GravityAccelerationComponents also matters

In [17]:

```
sns.boxplot(x='ActivityName', y='angleXgravityMean', data=train)
plt.axhline(y=0.08, xmin=0.1, xmax=0.9, c='m', dashes=(5,3))
plt.title('Angle between X-axis and Gravity_mean', fontsize=15)
plt.xticks(rotation = 40)
plt.show()
```




Observations:

- If $\text{angleX.gravityMean} > 0$ then Activity is Laying.
- We can classify all datapoints belonging to Laying activity with just a single if else statement.

In [18]:

```
sns.boxplot(x='ActivityName', y='angleYgravityMean', data = train, showfliers=False)
plt.title('Angle between Y-axis and Gravity_mean', fontsize=15)
plt.xticks(rotation = 40)
plt.axhline(y=-0.22, xmin=0.1, xmax=0.8, dashes=(5,3), c='m')
plt.show()
```



Apply t-sne on the data

In [46]:

```
import numpy as np
from sklearn.manifold import TSNE
import matplotlib.pyplot as plt
import seaborn as sns
```

In [47]:

```
# performs t-sne with different perplexity values and their repective plots..

def perform_tsne(X_data, y_data, perplexities, n_iter=1000, img_name_prefix='t-sne'):

    for index,perplexity in enumerate(perplexities):
        # perform t-sne
        print('\nperforming tsne with perplexity {} and with {} iterations at max'.format(perplexity, n_iter))
        X_reduced = TSNE(verbose=2, perplexity=perplexity).fit_transform(X_data)
        print('Done..')

        # prepare the data for seaborn
        print('Creating plot for this t-sne visualization..')
        df = pd.DataFrame({'x':X_reduced[:,0], 'y':X_reduced[:,1] , 'label':y_data})

        # draw the plot in appropriate place in the grid
        sns.lmplot(data=df, x='x', y='y', hue='label', fit_reg=False, size=8,\
                    palette="Set1",markers=['^','v','s','o', '1','2'])
        plt.title("perplexity : {} and max_iter : {}".format(perplexity, n_iter))
        img_name = img_name_prefix + '_perp_{}_iter_{}.png'.format(perplexity, n_iter)
        print('saving this plot as image in present working directory...')
        plt.savefig(img_name)
        plt.show()
        print('Done')
```


In [48]:

```
X_pre_tsne = train.drop(['subject', 'Activity','ActivityName'], axis=1)
y_pre_tsne = train['ActivityName']
perform_tsne(X_data = X_pre_tsne,y_data=y_pre_tsne, perplexities =[2,5,10,20,50])
```


```
performing tsne with perplexity 2 and with 1000 iterations at max
[t-SNE] Computing 7 nearest neighbors...
[t-SNE] Indexed 7352 samples in 0.426s...
[t-SNE] Computed neighbors for 7352 samples in 72.001s...
[t-SNE] Computed conditional probabilities for sample 1000 / 7352
[t-SNE] Computed conditional probabilities for sample 2000 / 7352
[t-SNE] Computed conditional probabilities for sample 3000 / 7352
[t-SNE] Computed conditional probabilities for sample 4000 / 7352
[t-SNE] Computed conditional probabilities for sample 5000 / 7352
[t-SNE] Computed conditional probabilities for sample 6000 / 7352
[t-SNE] Computed conditional probabilities for sample 7000 / 7352
[t-SNE] Computed conditional probabilities for sample 7352 / 7352
[t-SNE] Mean sigma: 0.635855
[t-SNE] Computed conditional probabilities in 0.071s
[t-SNE] Iteration 50: error = 124.8017578, gradient norm = 0.0253939 (50 i
terations in 16.625s)
[t-SNE] Iteration 100: error = 107.2019501, gradient norm = 0.0284782 (50
iterations in 9.735s)
[t-SNE] Iteration 150: error = 100.9872894, gradient norm = 0.0185151 (50
iterations in 5.346s)
[t-SNE] Iteration 200: error = 97.6054382, gradient norm = 0.0142084 (50 i
terations in 7.013s)
[t-SNE] Iteration 250: error = 95.3084183, gradient norm = 0.0132592 (50 i
terations in 5.703s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 95.308
418
[t-SNE] Iteration 300: error = 4.1209540, gradient norm = 0.0015668 (50 it
erations in 7.156s)
[t-SNE] Iteration 350: error = 3.2113254, gradient norm = 0.0009953 (50 it
erations in 8.022s)
[t-SNE] Iteration 400: error = 2.7819963, gradient norm = 0.0007203 (50 it
erations in 9.419s)
[t-SNE] Iteration 450: error = 2.5178111, gradient norm = 0.0005655 (50 it
erations in 9.370s)
[t-SNE] Iteration 500: error = 2.3341548, gradient norm = 0.0004804 (50 it
erations in 7.681s)
[t-SNE] Iteration 550: error = 2.1961622, gradient norm = 0.0004183 (50 it
erations in 7.097s)
[t-SNE] Iteration 600: error = 2.0867445, gradient norm = 0.0003664 (50 it
erations in 9.274s)
[t-SNE] Iteration 650: error = 1.9967778, gradient norm = 0.0003279 (50 it
erations in 7.697s)
[t-SNE] Iteration 700: error = 1.9210005, gradient norm = 0.0002984 (50 it
erations in 8.174s)
[t-SNE] Iteration 750: error = 1.8558111, gradient norm = 0.0002776 (50 it
erations in 9.747s)
[t-SNE] Iteration 800: error = 1.7989457, gradient norm = 0.0002569 (50 it
erations in 8.687s)
[t-SNE] Iteration 850: error = 1.7490212, gradient norm = 0.0002394 (50 it
erations in 8.407s)
[t-SNE] Iteration 900: error = 1.7043383, gradient norm = 0.0002224 (50 it
erations in 8.351s)
[t-SNE] Iteration 950: error = 1.6641431, gradient norm = 0.0002098 (50 it
erations in 7.841s)
[t-SNE] Iteration 1000: error = 1.6279151, gradient norm = 0.0001989 (50 i
terations in 5.623s)
[t-SNE] Error after 1000 iterations: 1.627915
Done..
Creating plot for this t-sne visualization..
saving this plot as image in present working directory...
```



Done

```
performing tsne with perplexity 5 and with 1000 iterations at max
[t-SNE] Computing 16 nearest neighbors...
[t-SNE] Indexed 7352 samples in 0.263s...
[t-SNE] Computed neighbors for 7352 samples in 48.983s...
[t-SNE] Computed conditional probabilities for sample 1000 / 7352
[t-SNE] Computed conditional probabilities for sample 2000 / 7352
[t-SNE] Computed conditional probabilities for sample 3000 / 7352
[t-SNE] Computed conditional probabilities for sample 4000 / 7352
[t-SNE] Computed conditional probabilities for sample 5000 / 7352
[t-SNE] Computed conditional probabilities for sample 6000 / 7352
[t-SNE] Computed conditional probabilities for sample 7000 / 7352
[t-SNE] Computed conditional probabilities for sample 7352 / 7352
[t-SNE] Mean sigma: 0.961265
[t-SNE] Computed conditional probabilities in 0.122s
[t-SNE] Iteration 50: error = 114.1862640, gradient norm = 0.0184120 (50 i
terations in 55.655s)
[t-SNE] Iteration 100: error = 97.6535568, gradient norm = 0.0174309 (50 i
terations in 12.580s)
[t-SNE] Iteration 150: error = 93.1900101, gradient norm = 0.0101048 (50 i
terations in 9.180s)
[t-SNE] Iteration 200: error = 91.2315445, gradient norm = 0.0074560 (50 i
terations in 10.340s)
[t-SNE] Iteration 250: error = 90.0714417, gradient norm = 0.0057667 (50 i
terations in 9.458s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 90.071
442
[t-SNE] Iteration 300: error = 3.5796804, gradient norm = 0.0014691 (50 it
erations in 8.718s)
[t-SNE] Iteration 350: error = 2.8173938, gradient norm = 0.0007508 (50 it
erations in 10.180s)
[t-SNE] Iteration 400: error = 2.4344938, gradient norm = 0.0005251 (50 it
erations in 10.506s)
[t-SNE] Iteration 450: error = 2.2156141, gradient norm = 0.0004069 (50 it
erations in 10.072s)
[t-SNE] Iteration 500: error = 2.0703306, gradient norm = 0.0003340 (50 it
erations in 10.511s)
[t-SNE] Iteration 550: error = 1.9646366, gradient norm = 0.0002816 (50 it
erations in 9.792s)
[t-SNE] Iteration 600: error = 1.8835558, gradient norm = 0.0002471 (50 it
erations in 9.098s)
[t-SNE] Iteration 650: error = 1.8184001, gradient norm = 0.0002184 (50 it
erations in 8.656s)
[t-SNE] Iteration 700: error = 1.7647167, gradient norm = 0.0001961 (50 it
erations in 9.063s)
[t-SNE] Iteration 750: error = 1.7193680, gradient norm = 0.0001796 (50 it
erations in 9.754s)
[t-SNE] Iteration 800: error = 1.6803776, gradient norm = 0.0001655 (50 it
erations in 9.540s)
[t-SNE] Iteration 850: error = 1.6465144, gradient norm = 0.0001538 (50 it
erations in 9.953s)
[t-SNE] Iteration 900: error = 1.6166563, gradient norm = 0.0001421 (50 it
erations in 10.270s)
[t-SNE] Iteration 950: error = 1.5901035, gradient norm = 0.0001335 (50 it
erations in 6.609s)
[t-SNE] Iteration 1000: error = 1.5664237, gradient norm = 0.0001257 (50 i
terations in 8.553s)
[t-SNE] Error after 1000 iterations: 1.566424
Done..
```


Creating plot for this t-sne visualization..
saving this plot as image in present working directory...



Done

```
performing tsne with perplexity 10 and with 1000 iterations at max
[t-SNE] Computing 31 nearest neighbors...
[t-SNE] Indexed 7352 samples in 0.410s...
[t-SNE] Computed neighbors for 7352 samples in 64.801s...
[t-SNE] Computed conditional probabilities for sample 1000 / 7352
[t-SNE] Computed conditional probabilities for sample 2000 / 7352
[t-SNE] Computed conditional probabilities for sample 3000 / 7352
[t-SNE] Computed conditional probabilities for sample 4000 / 7352
[t-SNE] Computed conditional probabilities for sample 5000 / 7352
[t-SNE] Computed conditional probabilities for sample 6000 / 7352
[t-SNE] Computed conditional probabilities for sample 7000 / 7352
[t-SNE] Computed conditional probabilities for sample 7352 / 7352
[t-SNE] Mean sigma: 1.133828
[t-SNE] Computed conditional probabilities in 0.214s
[t-SNE] Iteration 50: error = 106.0169220, gradient norm = 0.0194293 (50 i
terations in 24.550s)
[t-SNE] Iteration 100: error = 90.3036194, gradient norm = 0.0097653 (50 i
terations in 11.936s)
[t-SNE] Iteration 150: error = 87.3132935, gradient norm = 0.0053059 (50 i
terations in 11.246s)
[t-SNE] Iteration 200: error = 86.1169128, gradient norm = 0.0035844 (50 i
terations in 11.864s)
[t-SNE] Iteration 250: error = 85.4133606, gradient norm = 0.0029100 (50 i
terations in 11.944s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 85.413
361
[t-SNE] Iteration 300: error = 3.1394315, gradient norm = 0.0013976 (50 it
erations in 11.742s)
[t-SNE] Iteration 350: error = 2.4929206, gradient norm = 0.0006466 (50 it
erations in 11.627s)
[t-SNE] Iteration 400: error = 2.1733041, gradient norm = 0.0004230 (50 it
erations in 11.846s)
[t-SNE] Iteration 450: error = 1.9884514, gradient norm = 0.0003124 (50 it
erations in 11.405s)
[t-SNE] Iteration 500: error = 1.8702440, gradient norm = 0.0002514 (50 it
erations in 11.320s)
[t-SNE] Iteration 550: error = 1.7870129, gradient norm = 0.0002107 (50 it
erations in 12.009s)
[t-SNE] Iteration 600: error = 1.7246909, gradient norm = 0.0001824 (50 it
erations in 10.632s)
[t-SNE] Iteration 650: error = 1.6758548, gradient norm = 0.0001590 (50 it
erations in 11.270s)
[t-SNE] Iteration 700: error = 1.6361949, gradient norm = 0.0001451 (50 it
erations in 12.072s)
[t-SNE] Iteration 750: error = 1.6034756, gradient norm = 0.0001305 (50 it
erations in 11.607s)
[t-SNE] Iteration 800: error = 1.5761518, gradient norm = 0.0001188 (50 it
erations in 9.409s)
[t-SNE] Iteration 850: error = 1.5527289, gradient norm = 0.0001113 (50 it
erations in 8.309s)
[t-SNE] Iteration 900: error = 1.5328671, gradient norm = 0.0001021 (50 it
erations in 9.433s)
[t-SNE] Iteration 950: error = 1.5152045, gradient norm = 0.0000974 (50 it
erations in 11.488s)
[t-SNE] Iteration 1000: error = 1.4999681, gradient norm = 0.0000933 (50 i
terations in 10.593s)
[t-SNE] Error after 1000 iterations: 1.499968
Done..
```


Creating plot for this t-sne visualization..
saving this plot as image in present working directory...



Done

```
performing tsne with perplexity 20 and with 1000 iterations at max
[t-SNE] Computing 61 nearest neighbors...
[t-SNE] Indexed 7352 samples in 0.425s...
[t-SNE] Computed neighbors for 7352 samples in 61.792s...
[t-SNE] Computed conditional probabilities for sample 1000 / 7352
[t-SNE] Computed conditional probabilities for sample 2000 / 7352
[t-SNE] Computed conditional probabilities for sample 3000 / 7352
[t-SNE] Computed conditional probabilities for sample 4000 / 7352
[t-SNE] Computed conditional probabilities for sample 5000 / 7352
[t-SNE] Computed conditional probabilities for sample 6000 / 7352
[t-SNE] Computed conditional probabilities for sample 7000 / 7352
[t-SNE] Computed conditional probabilities for sample 7352 / 7352
[t-SNE] Mean sigma: 1.274335
[t-SNE] Computed conditional probabilities in 0.355s
[t-SNE] Iteration 50: error = 97.5202179, gradient norm = 0.0223863 (50 it
erations in 21.168s)
[t-SNE] Iteration 100: error = 83.9500732, gradient norm = 0.0059110 (50 i
terations in 17.306s)
[t-SNE] Iteration 150: error = 81.8804779, gradient norm = 0.0035797 (50 i
terations in 14.258s)
[t-SNE] Iteration 200: error = 81.1615143, gradient norm = 0.0022536 (50 i
terations in 14.130s)
[t-SNE] Iteration 250: error = 80.7704086, gradient norm = 0.0018108 (50 i
terations in 15.340s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 80.770
409
[t-SNE] Iteration 300: error = 2.6957574, gradient norm = 0.0012993 (50 it
erations in 13.605s)
[t-SNE] Iteration 350: error = 2.1637220, gradient norm = 0.0005765 (50 it
erations in 13.248s)
[t-SNE] Iteration 400: error = 1.9143614, gradient norm = 0.0003474 (50 it
erations in 14.774s)
[t-SNE] Iteration 450: error = 1.7684202, gradient norm = 0.0002458 (50 it
erations in 15.502s)
[t-SNE] Iteration 500: error = 1.6744757, gradient norm = 0.0001923 (50 it
erations in 14.808s)
[t-SNE] Iteration 550: error = 1.6101606, gradient norm = 0.0001575 (50 it
erations in 14.043s)
[t-SNE] Iteration 600: error = 1.5641028, gradient norm = 0.0001344 (50 it
erations in 15.769s)
[t-SNE] Iteration 650: error = 1.5291905, gradient norm = 0.0001182 (50 it
erations in 15.834s)
[t-SNE] Iteration 700: error = 1.5024391, gradient norm = 0.0001055 (50 it
erations in 15.398s)
[t-SNE] Iteration 750: error = 1.4809053, gradient norm = 0.0000965 (50 it
erations in 14.594s)
[t-SNE] Iteration 800: error = 1.4631859, gradient norm = 0.0000884 (50 it
erations in 15.025s)
[t-SNE] Iteration 850: error = 1.4486470, gradient norm = 0.0000832 (50 it
erations in 14.060s)
[t-SNE] Iteration 900: error = 1.4367288, gradient norm = 0.0000804 (50 it
erations in 12.389s)
[t-SNE] Iteration 950: error = 1.4270191, gradient norm = 0.0000761 (50 it
erations in 10.392s)
[t-SNE] Iteration 1000: error = 1.4189968, gradient norm = 0.0000787 (50 i
terations in 12.355s)
[t-SNE] Error after 1000 iterations: 1.418997
Done..
```



Creating plot for this t-sne visualization..
saving this plot as image in present working directory...



Done

```
performing tsne with perplexity 50 and with 1000 iterations at max
[t-SNE] Computing 151 nearest neighbors...
[t-SNE] Indexed 7352 samples in 0.376s...
[t-SNE] Computed neighbors for 7352 samples in 73.164s...
[t-SNE] Computed conditional probabilities for sample 1000 / 7352
[t-SNE] Computed conditional probabilities for sample 2000 / 7352
[t-SNE] Computed conditional probabilities for sample 3000 / 7352
[t-SNE] Computed conditional probabilities for sample 4000 / 7352
[t-SNE] Computed conditional probabilities for sample 5000 / 7352
[t-SNE] Computed conditional probabilities for sample 6000 / 7352
[t-SNE] Computed conditional probabilities for sample 7000 / 7352
[t-SNE] Computed conditional probabilities for sample 7352 / 7352
[t-SNE] Mean sigma: 1.437672
[t-SNE] Computed conditional probabilities in 0.844s
[t-SNE] Iteration 50: error = 86.1525574, gradient norm = 0.0242986 (50 it
erations in 36.249s)
[t-SNE] Iteration 100: error = 75.9874649, gradient norm = 0.0061005 (50 i
terations in 30.453s)
[t-SNE] Iteration 150: error = 74.7072296, gradient norm = 0.0024708 (50 i
terations in 28.461s)
[t-SNE] Iteration 200: error = 74.2736282, gradient norm = 0.0018644 (50 i
terations in 27.735s)
[t-SNE] Iteration 250: error = 74.0722427, gradient norm = 0.0014078 (50 i
terations in 26.835s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 74.072
243
[t-SNE] Iteration 300: error = 2.1539080, gradient norm = 0.0011796 (50 it
erations in 25.445s)
[t-SNE] Iteration 350: error = 1.7567128, gradient norm = 0.0004845 (50 it
erations in 21.282s)
[t-SNE] Iteration 400: error = 1.5888531, gradient norm = 0.0002798 (50 it
erations in 21.015s)
[t-SNE] Iteration 450: error = 1.4956820, gradient norm = 0.0001894 (50 it
erations in 23.332s)
[t-SNE] Iteration 500: error = 1.4359720, gradient norm = 0.0001420 (50 it
erations in 23.083s)
[t-SNE] Iteration 550: error = 1.3947564, gradient norm = 0.0001117 (50 it
erations in 19.626s)
[t-SNE] Iteration 600: error = 1.3653858, gradient norm = 0.0000949 (50 it
erations in 22.752s)
[t-SNE] Iteration 650: error = 1.3441534, gradient norm = 0.0000814 (50 it
erations in 23.972s)
[t-SNE] Iteration 700: error = 1.3284039, gradient norm = 0.0000742 (50 it
erations in 20.636s)
[t-SNE] Iteration 750: error = 1.3171139, gradient norm = 0.0000700 (50 it
erations in 20.407s)
[t-SNE] Iteration 800: error = 1.3085558, gradient norm = 0.0000657 (50 it
erations in 24.951s)
[t-SNE] Iteration 850: error = 1.3017821, gradient norm = 0.0000603 (50 it
erations in 24.719s)
[t-SNE] Iteration 900: error = 1.2962619, gradient norm = 0.0000586 (50 it
erations in 24.500s)
[t-SNE] Iteration 950: error = 1.2914882, gradient norm = 0.0000573 (50 it
erations in 24.132s)
[t-SNE] Iteration 1000: error = 1.2874244, gradient norm = 0.0000546 (50 i
terations in 22.840s)
[t-SNE] Error after 1000 iterations: 1.287424
Done..
```

Creating plot for this t-sne visualization..
saving this plot as image in present working directory...



Done

Applying traditional ML models on the handcrafted features

In [1]:

```
import numpy as np
import pandas as pd
```

Obtain the train and test data

In [2]:

```
train = pd.read_csv('UCI_HAR_dataset/csv_files/train.csv')
test = pd.read_csv('UCI_HAR_dataset/csv_files/test.csv')
print(train.shape, test.shape)
```

(7352, 564) (2947, 564)

In [3]:

```
train.head(3)
```

Out[3]:

	tBodyAccmeanX	tBodyAccmeanY	tBodyAccmeanZ	tBodyAccstdX	tBodyAccstdY
0	0.288585	-0.020294	-0.132905	-0.995279	-0.983111
1	0.278419	-0.016411	-0.123520	-0.998245	-0.975300
2	0.279653	-0.019467	-0.113462	-0.995380	-0.967187

3 rows × 564 columns

In [4]:

```
# get X_train and y_train from csv files
X_train = train.drop(['subject', 'Activity', 'ActivityName'], axis=1)
y_train = train.ActivityName
```

In [5]:

```
# get X_test and y_test from test csv file
X_test = test.drop(['subject', 'Activity', 'ActivityName'], axis=1)
y_test = test.ActivityName
```

In [6]:

```
print('X_train and y_train : ({},{})'.format(X_train.shape, y_train.shape))
print('X_test and y_test : ({},{})'.format(X_test.shape, y_test.shape))
```

X_train and y_train : ((7352, 561),(7352,))

X_test and y_test : ((2947, 561),(2947,))

Let's model with our data

Labels that are useful in plotting confusion matrix

In [7]:

```
labels=['LAYING', 'SITTING', 'STANDING', 'WALKING', 'WALKING_DOWNSTAIRS', 'WALKING_UPSTAIRS']
```

Function to plot the confusion matrix

In [8]:

```
import itertools
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix
plt.rcParams["font.family"] = 'DejaVu Sans'

def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=90)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
```

Generic function to run any model specified

In [9]:

```

from datetime import datetime
def perform_model(model, X_train, y_train, X_test, y_test, class_labels, cm_normalize=True, \
                  print_cm=True, cm_cmap=plt.cm.Greens):

    # to store results at various phases
    results = dict()

    # time at which model starts training
    train_start_time = datetime.now()
    print('training the model..')
    model.fit(X_train, y_train)
    print('Done \n \n')
    train_end_time = datetime.now()
    results['training_time'] = train_end_time - train_start_time
    print('training_time(HH:MM:SS.ms) - {}'.format(results['training_time']))

    # predict test data
    print('Predicting test data')
    test_start_time = datetime.now()
    y_pred = model.predict(X_test)
    test_end_time = datetime.now()
    print('Done \n \n')
    results['testing_time'] = test_end_time - test_start_time
    print('testing time(HH:MM:SS.ms) - {}'.format(results['testing_time']))
    results['predicted'] = y_pred

    # calculate overall accuracy of the model
    accuracy = metrics.accuracy_score(y_true=y_test, y_pred=y_pred)
    # store accuracy in results
    results['accuracy'] = accuracy
    print('-----')
    print('|          Accuracy          |')
    print('-----')
    print('\n    {}'.format(accuracy))

    # confusion matrix
    cm = metrics.confusion_matrix(y_test, y_pred)
    results['confusion_matrix'] = cm
    if print_cm:
        print('-----')
        print('| Confusion Matrix |')
        print('-----')
        print('\n {}'.format(cm))

    # plot confusion matrix
    plt.figure(figsize=(8,8))
    plt.grid(b=False)
    plot_confusion_matrix(cm, classes=class_labels, normalize=True, title='Normalized c
onfusion matrix', cmap = cm_cmap)
    plt.show()

    # get classification report
    print('-----')
    print('| Classification Report |')

```

```

print('-----')
classification_report = metrics.classification_report(y_test, y_pred)
# store report in results
results['classification_report'] = classification_report
print(classification_report)

# add the trained model to the results
results['model'] = model

return results

```

Method to print the gridsearch Attributes

In [10]:

```

def print_grid_search_attributes(model):
    # Estimator that gave highest score among all the estimators formed in GridSearch
    print('-----')
    print('|          Best Estimator          |')
    print('-----')
    print('\n\t{}\n'.format(model.best_estimator_))

    # parameters that gave best results while performing grid search
    print('-----')
    print('|          Best parameters          |')
    print('-----')
    print('\tParameters of best estimator : \n\n\t{}\n'.format(model.best_params_))

    # number of cross validation splits
    print('-----')
    print('|    No of CrossValidation sets    |')
    print('-----')
    print('\n\tTotal numbre of cross validation sets: {}\n'.format(model.n_splits_))

    # Average cross validated score of the best estimator, from the Grid Search
    print('-----')
    print('|          Best Score          |')
    print('-----')
    print('\n\tAverage Cross Validate scores of best estimator : \n\n\t{}\n'.format(model.best_score_))

```

1. Logistic Regression with Grid Search

In [11]:

```
from sklearn import linear_model  
from sklearn import metrics  
  
from sklearn.model_selection import GridSearchCV
```


In [12]:

```
# start Grid search
parameters = {'C':[0.01, 0.1, 1, 10, 20, 30], 'penalty':['l2','l1']}
log_reg = linear_model.LogisticRegression()
log_reg_grid = GridSearchCV(log_reg, param_grid=parameters, cv=3, verbose=1, n_jobs=-1)
log_reg_grid_results = perform_model(log_reg_grid, X_train, y_train, X_test, y_test, c
lass_labels=labels)
```

training the model..

Fitting 3 folds for each of 12 candidates, totalling 36 fits

[Parallel(n_jobs=-1)]: Done 36 out of 36 | elapsed: 1.2min finished

Done

training_time(HH:MM:SS.ms) - 0:01:25.843810

Predicting test data

Done


testing time(HH:MM:SS.ms) - 0:00:00.009192

Accuracy

0.9626739056667798

Confusion Matrix

```
[[537  0  0  0  0  0]
 [ 1428 58  0  0  4]
 [  0 12 519  1  0  0]
 [  0  0  0 495  1  0]
 [  0  0  0  3 409  8]
 [  0  0  0 22  0 449]]
```




Classification Report

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.97	0.87	0.92	491
STANDING	0.90	0.98	0.94	532
WALKING	0.95	1.00	0.97	496
WALKING_DOWNSTAIRS	1.00	0.97	0.99	420
WALKING_UPSTAIRS	0.97	0.95	0.96	471
avg / total	0.96	0.96	0.96	2947

In [13]:

```
plt.figure(figsize=(8,8))
plt.grid(b=False)
plot_confusion_matrix(log_reg_grid_results['confusion_matrix'], classes=labels, cmap=plt.cm.Greens, )
plt.show()
```



In [14]:

```
# observe the attributes of the model
print_grid_search_attributes(log_reg_grid_results['model'])
```

```
-----
|      Best Estimator      |
|-----|
```

```
LogisticRegression(C=30, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                    verbose=0, warm_start=False)
```

```
-----
|    Best parameters      |
|-----|
```

Parameters of best estimator :

```
{'C': 30, 'penalty': 'l2'}
```

```
-----
| No of CrossValidation sets |
|-----|
```

Total number of cross validation sets: 3

```
-----
|      Best Score        |
|-----|
```

Average Cross Validate scores of best estimator :

```
0.9461371055495104
```

2. Linear SVC with GridSearch

In [15]:

```
from sklearn.svm import LinearSVC
```

In [16]:

```
parameters = {'C':[0.125, 0.5, 1, 2, 8, 16]}  
lr_svc = LinearSVC(tol=0.00005)  
lr_svc_grid = GridSearchCV(lr_svc, param_grid=parameters, n_jobs=-1, verbose=1)  
lr_svc_grid_results = perform_model(lr_svc_grid, X_train, y_train, X_test, y_test, clas  
s_labels=labels)
```

training the model..

Fitting 3 folds for each of 6 candidates, totalling 18 fits

[Parallel(n_jobs=-1)]: Done 18 out of 18 | elapsed: 24.9s finished

Done

training_time(HH:MM:SS.ms) - 0:00:32.951942

Predicting test data

Done


testing time(HH:MM:SS.ms) - 0:00:00.012182

```
-----  
| Accuracy |  
-----
```

0.9660671869697998

```
-----  
| Confusion Matrix |  
-----
```

```
[[537  0  0  0  0  0]  
[ 2 426 58  0  0  5]  
[ 0 14 518  0  0  0]  
[ 0  0  0 495  0  1]  
[ 0  0  0  2 413  5]  
[ 0  0  0 12  1 458]]
```




```
-----
| Classification Report |
-----
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.97	0.87	0.92	491
STANDING	0.90	0.97	0.94	532
WALKING	0.97	1.00	0.99	496
WALKING_DOWNSTAIRS	1.00	0.98	0.99	420
WALKING_UPSTAIRS	0.98	0.97	0.97	471
avg / total	0.97	0.97	0.97	2947

In [17]:

```
print_grid_search_attributes(lr_svc_grid_results['model'])
```

```
-----
| Best Estimator |
-----
```

```
LinearSVC(C=8, class_weight=None, dual=True, fit_intercept=True,
intercept_scaling=1, loss='squared_hinge', max_iter=1000,
multi_class='ovr', penalty='l2', random_state=None, tol=5e-05,
verbose=0)
```

```
-----
| Best parameters |
-----
```

Parameters of best estimator :

```
{'C': 8}
```

```
-----
| No of CrossValidation sets |
-----
```

Total numbre of cross validation sets: 3

```
-----
| Best Score |
-----
```

Average Cross Validate scores of best estimator :

```
0.9465451577801959
```

3. Kernel SVM with GridSearch

In [18]:

```
from sklearn.svm import SVC
parameters = {'C':[2,8,16],\
              'gamma': [ 0.0078125, 0.125, 2]}
rbf_svm = SVC(kernel='rbf')
rbf_svm_grid = GridSearchCV(rbf_svm,param_grid=parameters, n_jobs=-1)
rbf_svm_grid_results = perform_model(rbf_svm_grid, X_train, y_train, X_test, y_test, class_labels=labels)
```

training the model..
Done

training_time(HH:MM:SS.ms) - 0:05:46.182889

Predicting test data
Done


testing time(HH:MM:SS.ms) - 0:00:05.221285

Accuracy

0.9626739056667798

Confusion Matrix

```
[[537  0  0  0  0  0]
 [ 0 441 48  0  0  2]
 [ 0 12 520  0  0  0]
 [ 0  0  0 489  2  5]
 [ 0  0  0  4 397 19]
 [ 0  0  0 17  1 453]]
```



```

-----
| Classification Report |
-----

```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.97	0.90	0.93	491
STANDING	0.92	0.98	0.95	532
WALKING	0.96	0.99	0.97	496
WALKING_DOWNSTAIRS	0.99	0.95	0.97	420
WALKING_UPSTAIRS	0.95	0.96	0.95	471
avg / total	0.96	0.96	0.96	2947

In [19]:

```
print_grid_search_attributes(rbf_svm_grid_results['model'])
```

```

-----
| Best Estimator |
-----

```

```
SVC(C=16, cache_size=200, class_weight=None, coef0=0.0,
decision_function_shape='ovr', degree=3, gamma=0.0078125, kernel='rbf',
max_iter=-1, probability=False, random_state=None, shrinking=True,
tol=0.001, verbose=False)
```

```

-----
| Best parameters |
-----

```

Parameters of best estimator :

```
{'C': 16, 'gamma': 0.0078125}
```

```

-----
| No of CrossValidation sets |
-----

```

Total nombre of cross validation sets: 3

```

-----
| Best Score |
-----

```

Average Cross Validate scores of best estimator :

```
0.9440968443960827
```

4. Decision Trees with GridSearchCV

In [20]:

```
from sklearn.tree import DecisionTreeClassifier
parameters = {'max_depth':np.arange(3,10,2)}
dt = DecisionTreeClassifier()
dt_grid = GridSearchCV(dt,param_grid=parameters, n_jobs=-1)
dt_grid_results = perform_model(dt_grid, X_train, y_train, X_test, y_test, class_labels
=labels)
print_grid_search_attributes(dt_grid_results['model'])
```

training the model..
Done

training_time(HH:MM:SS.ms) - 0:00:19.476858

Predicting test data
Done


testing time(HH:MM:SS.ms) - 0:00:00.012858

Accuracy

0.8642687478791992

Confusion Matrix

```
[[537  0  0  0  0  0]
 [ 0 386 105  0  0  0]
 [ 0  93 439  0  0  0]
 [ 0  0  0 472 16  8]
 [ 0  0  0 15 344 61]
 [ 0  0  0 73 29 369]]
```




```

-----
| Classification Report |
-----

```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.81	0.79	0.80	491
STANDING	0.81	0.83	0.82	532
WALKING	0.84	0.95	0.89	496
WALKING_DOWNSTAIRS	0.88	0.82	0.85	420
WALKING_UPSTAIRS	0.84	0.78	0.81	471
avg / total	0.86	0.86	0.86	2947

```

-----
| Best Estimator |
-----

```

```

DecisionTreeClassifier(class_weight=None, criterion='gini', max_de
pth=7,
    max_features=None, max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, presort=False, random_state=Non
e,
    splitter='best')

```

```

-----
| Best parameters |
-----

```

Parameters of best estimator :

```
{'max_depth': 7}
```

```

-----
| No of CrossValidation sets |
-----

```

Total nombre of cross validation sets: 3

```

-----
| Best Score |
-----

```

Average Cross Validate scores of best estimator :

```
0.8369151251360174
```

5. Random Forest Classifier with GridSearch

In [21]:

```
from sklearn.ensemble import RandomForestClassifier
params = {'n_estimators': np.arange(10,201,20), 'max_depth':np.arange(3,15,2)}
rfc = RandomForestClassifier()
rfc_grid = GridSearchCV(rfc, param_grid=params, n_jobs=-1)
rfc_grid_results = perform_model(rfc_grid, X_train, y_train, X_test, y_test, class_labels=labels)
print_grid_search_attributes(rfc_grid_results['model'])
```

training the model..
Done

training_time(HH:MM:SS.ms) - 0:06:22.775270

Predicting test data
Done


testing time(HH:MM:SS.ms) - 0:00:00.025937

Accuracy

0.9131319986426875

Confusion Matrix

```
[[537  0  0  0  0  0]
 [ 0 427 64  0  0  0]
 [ 0 52 480  0  0  0]
 [ 0  0  0 484 10  2]
 [ 0  0  0 38 332 50]
 [ 0  0  0 34  6 431]]
```



Classification Report				

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.89	0.87	0.88	491
STANDING	0.88	0.90	0.89	532
WALKING	0.87	0.98	0.92	496
WALKING_DOWNSTAIRS	0.95	0.79	0.86	420
WALKING_UPSTAIRS	0.89	0.92	0.90	471
avg / total	0.92	0.91	0.91	2947

```
-----
| Best Estimator |
-----
```

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterio
n='gini',
                        max_depth=7, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=70, n_jobs=1,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
```

```
-----
| Best parameters |
-----
```

Parameters of best estimator :

```
{'max_depth': 7, 'n_estimators': 70}
```

```
-----
| No of CrossValidation sets |
-----
```

Total nombre of cross validation sets: 3

```
-----
| Best Score |
-----
```

Average Cross Validate scores of best estimator :

```
0.9141730141458106
```

6. Gradient Boosted Decision Trees With GridSearch

In [22]:

```
from sklearn.ensemble import GradientBoostingClassifier
param_grid = {'max_depth': np.arange(5,8,1), \
              'n_estimators':np.arange(130,170,10)}
gbdt = GradientBoostingClassifier()
gbdt_grid = GridSearchCV(gbdt, param_grid=param_grid, n_jobs=-1)
gbdt_grid_results = perform_model(gbdt_grid, X_train, y_train, X_test, y_test, class_labels=labels)
print_grid_search_attributes(gbdt_grid_results['model'])
```

```
training the model..
```

```
Done
```

```
training_time(HH:MM:SS.ms) - 0:28:03.653432
```

```
Predicting test data
```

```
Done
```


```
testing time(HH:MM:SS.ms) - 0:00:00.058843
```

```
-----  
|      Accuracy      |  
-----
```

```
0.9222938581608415
```

```
-----  
| Confusion Matrix |  
-----
```

```
[[537  0  0  0  0  0]  
[ 0 396 93  0  0  2]  
[ 0 37 495  0  0  0]  
[ 0  0  0 483  7  6]  
[ 0  0  0 10 374 36]  
[ 0  1  0 31  6 433]]
```




```

-----
| Classification Report |
-----

```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	537
SITTING	0.91	0.81	0.86	491
STANDING	0.84	0.93	0.88	532
WALKING	0.92	0.97	0.95	496
WALKING_DOWNSTAIRS	0.97	0.89	0.93	420
WALKING_UPSTAIRS	0.91	0.92	0.91	471
avg / total	0.92	0.92	0.92	2947

```

-----
| Best Estimator |
-----

```

```

GradientBoostingClassifier(criterion='friedman_mse', init=None,
    learning_rate=0.1, loss='deviance', max_depth=5,
    max_features=None, max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=140,
    presort='auto', random_state=None, subsample=1.0, verbose=0,
    warm_start=False)

```

```

-----
| Best parameters |
-----

```

Parameters of best estimator :

```
{'max_depth': 5, 'n_estimators': 140}
```

```

-----
| No of CrossValidation sets |
-----

```

Total nombre of cross validation sets: 3

```

-----
| Best Score |
-----

```

Average Cross Validate scores of best estimator :

```
0.904379760609358
```

7. Comparing all models

In [23]:

```
print('\n
Accuracy      Error')
print('
-----      -')
print('Logistic Regression : {:.04}%      {:.04}%'.format(log_reg_grid_results['accuracy'] * 100, \
100-(log_reg_grid_results['accuracy'] * 100)))

print('Linear SVC      : {:.04}%      {:.04}% '.format(lr_svc_grid_results['accuracy'] * 100, \
100-(lr_svc_grid_results['accuracy'] * 100)))

print('rbf SVM classifier : {:.04}%      {:.04}% '.format(rbf_svm_grid_results['accuracy'] * 100, \
100-(rbf_svm_grid_results['accuracy'] * 100)))

print('DecisionTree      : {:.04}%      {:.04}% '.format(dt_grid_results['accuracy'] * 100, \
100-(dt_grid_results['accuracy'] * 100)))

print('Random Forest      : {:.04}%      {:.04}% '.format(rfc_grid_results['accuracy'] * 100, \
100-(rfc_grid_results['accuracy'] * 100)))

print('GradientBoosting DT : {:.04}%      {:.04}% '.format(rfc_grid_results['accuracy'] * 100, \
100-(rfc_grid_results['accuracy'] * 100)))
```

	Accuracy -----	Error -----
Logistic Regression	: 96.27%	3.733%
Linear SVC	: 96.61%	3.393%
rbf SVM classifier	: 96.27%	3.733%
DecisionTree	: 86.43%	13.57%
Random Forest	: 91.31%	8.687%
GradientBoosting DT	: 91.31%	8.687%

We can choose **Logistic regression** or **Linear SVC** or **rbf SVM**.

Conclusion :

In the real world, domain-knowledge, EDA and feature-engineering matter most.

Human activity recognition using LSTM

In []:

```
# Importing Libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from time import time
from datetime import datetime
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Dropout
from keras.wrappers.scikit_learn import KerasClassifier
from keras.constraints import maxnorm
from keras.models import Sequential
from keras.layers import LSTM
from keras.layers.core import Dense, Dropout
from keras.layers.normalization import BatchNormalization
```

In [0]:

```
# Activities are the class labels
# It is a 6 class classification
ACTIVITIES = {
    0: 'WALKING',
    1: 'WALKING_UPSTAIRS',
    2: 'WALKING_DOWNSTAIRS',
    3: 'SITTING',
    4: 'STANDING',
    5: 'LAYING',
}

# Utility function to print the confusion matrix
def confusion_matrix(Y_true, Y_pred):
    Y_true = pd.Series([ACTIVITIES[y] for y in np.argmax(Y_true, axis=1)])
    Y_pred = pd.Series([ACTIVITIES[y] for y in np.argmax(Y_pred, axis=1)])

    return pd.crosstab(Y_true, Y_pred, rownames=['True'], colnames=['Pred'])
```

1.0 Data

In [0]:

```

# Raw data signals
# Signals are from Accelerometer and Gyroscope
# The signals are in x,y,z directions
# Sensor signals are filtered to have only body acceleration
# excluding the acceleration due to gravity
# Triaxial acceleration from the accelerometer is total acceleration
SIGNALS = [
    "body_acc_x",
    "body_acc_y",
    "body_acc_z",
    "body_gyro_x",
    "body_gyro_y",
    "body_gyro_z",
    "total_acc_x",
    "total_acc_y",
    "total_acc_z"
]

```

In [0]:

```

# Utility function to read the data from csv file
def _read_csv(filename):
    return pd.read_csv(filename, delim_whitespace=True, header=None)

# Utility function to load the load
def load_signals(subset):
    signals_data = []

    for signal in SIGNALS:
        filename = f'UCI_HAR_Dataset/{subset}/Inertial Signals/{signal}_{subset}.txt'
        signals_data.append(
            _read_csv(filename).as_matrix()
        )

    # Transpose is used to change the dimensionality of the output,
    # aggregating the signals by combination of sample/timestep.
    # Resultant shape is (7352 train/2947 test samples, 128 timesteps, 9 signals)
    return np.transpose(signals_data, (1, 2, 0))

```

In [0]:

```

def load_y(subset):
    """
    The objective that we are trying to predict is a integer, from 1 to 6,
    that represents a human activity. We return a binary representation of
    every sample objective as a 6 bits vector using One Hot Encoding
    (https://pandas.pydata.org/pandas-docs/stable/generated/pandas.get\_dummies.html)
    """
    filename = f'UCI_HAR_Dataset/{subset}/y_{subset}.txt'
    y = _read_csv(filename)[0]

    return pd.get_dummies(y).as_matrix()

```

In [0]:

```
def load_data():  
    """  
    Obtain the dataset from multiple files.  
    Returns: X_train, X_test, y_train, y_test  
    """  
    X_train, X_test = load_signals('train'), load_signals('test')  
    y_train, y_test = load_y('train'), load_y('test')  
  
    return X_train, X_test, y_train, y_test
```

In [0]:

```
# Importing tensorflow  
np.random.seed(42)  
import tensorflow as tf  
tf.random.set_seed(42)
```

In [0]:

```
# Configuring a session  
session_conf = tf.compat.v1.ConfigProto(  
    intra_op_parallelism_threads=1,  
    inter_op_parallelism_threads=1  
)
```

In [0]:

```
# Import Keras  
from keras import backend as K  
sess = tf.Session(graph=tf.get_default_graph(), config=session_conf)  
K.set_session(sess)
```

In [0]:

```
# Utility function to count the number of classes  
def _count_classes(y):  
    return len(set([tuple(category) for category in y]))
```

In [0]:

```
# Loading the train and test data  
import warnings  
warnings.filterwarnings("ignore")  
  
X_train, X_test, Y_train, Y_test = load_data()
```

In [0]:

```
type(X_train)
```

Out[0]:

```
numpy.ndarray
```

In [0]:

```
print((X_train[0][0]))
```

```
[ 1.808515e-04  1.076681e-02  5.556068e-02  3.019122e-02  6.601362e-02  
 2.285864e-02  1.012817e+00 -1.232167e-01  1.029341e-01]
```

In [0]:

```
print((X_train[0]))
```

```
[[ 1.808515e-04  1.076681e-02  5.556068e-02 ...  1.012817e+00  
 -1.232167e-01  1.029341e-01]  
 [ 1.013856e-02  6.579480e-03  5.512483e-02 ...  1.022833e+00  
 -1.268756e-01  1.056872e-01]  
 [ 9.275574e-03  8.928878e-03  4.840473e-02 ...  1.022028e+00  
 -1.240037e-01  1.021025e-01]  
 ...  
 [-1.147484e-03  1.714439e-04  2.647864e-03 ...  1.018445e+00  
 -1.240696e-01  1.003852e-01]  
 [-2.222655e-04  1.574181e-03  2.381057e-03 ...  1.019372e+00  
 -1.227451e-01  9.987355e-02]  
 [ 1.575500e-03  3.070189e-03 -2.269757e-03 ...  1.021171e+00  
 -1.213260e-01  9.498741e-02]]
```

In [0]:

```
timesteps = len(X_train[0])  
input_dim = len(X_train[0][0])  
n_classes = _count_classes(Y_train)  
  
print(timesteps)  
print(input_dim)  
print(len(X_train))
```

```
128  
9  
7352
```

In [0]:

```
print(n_classes)
```

```
6
```

In [0]:

```
np.save('X_train', X_train)  
np.save('X_test', X_test)  
np.save('Y_train', Y_train)  
np.save('Y_test', Y_test)
```

In [3]:

```
from zipfile import ZipFile
file_name="/content/Colab.zip"

with ZipFile(file_name,'r') as zip:
    zip.extractall()
    print('Done')
```

Done

In [0]:

```
X_train= np.load('/content/Colab/X_train.npy')
X_test= np.load('/content/Colab/X_test.npy')
Y_train= np.load('/content/Colab/Y_train.npy')
Y_test= np.load('/content/Colab/Y_test.npy')
```

In [0]:

```
Y_test= np.load('/content/Colab/Y_test.npy')
```

2.0 Simple base model without hyperparameter tuning

In [0]:

```
# Initializing parameters
epochs = 30
batch_size = 16
n_hidden = 32
```

In [0]:

```
# Initiliazing the sequential model
model = Sequential()
# Configuring the parameters
model.add(LSTM(n_hidden, input_shape=(timesteps, input_dim)))
# Adding a dropout layer
model.add(Dropout(0.5))
# Adding a dense output layer with sigmoid activation
model.add(Dense(n_classes, activation='sigmoid'))
model.summary()
```

Layer (type)	Output Shape	Param #
=====		
lstm_3 (LSTM)	(None, 32)	5376

dropout_3 (Dropout)	(None, 32)	0

dense_3 (Dense)	(None, 6)	198
=====		
Total params: 5,574		
Trainable params: 5,574		
Non-trainable params: 0		

In [0]:

```
# Compiling the model
model.compile(loss='categorical_crossentropy',
              optimizer='rmsprop',
              metrics=['accuracy'])
```


In [0]:

```
# Training the model
model.fit(X_train,
          Y_train,
          batch_size=batch_size,
          validation_data=(X_test, Y_test),
          epochs=epochs)
```

Train on 7352 samples, validate on 2947 samples

Epoch 1/30

7352/7352 [=====] - 92s 13ms/step - loss: 1.3018
- acc: 0.4395 - val_loss: 1.1254 - val_acc: 0.4662

Epoch 2/30

7352/7352 [=====] - 94s 13ms/step - loss: 0.9666
- acc: 0.5880 - val_loss: 0.9491 - val_acc: 0.5714

Epoch 3/30

7352/7352 [=====] - 97s 13ms/step - loss: 0.7812
- acc: 0.6408 - val_loss: 0.8286 - val_acc: 0.5850

Epoch 4/30

7352/7352 [=====] - 95s 13ms/step - loss: 0.6941
- acc: 0.6574 - val_loss: 0.7297 - val_acc: 0.6128

Epoch 5/30

7352/7352 [=====] - 92s 13ms/step - loss: 0.6336
- acc: 0.6912 - val_loss: 0.7359 - val_acc: 0.6787

Epoch 6/30

7352/7352 [=====] - 94s 13ms/step - loss: 0.5859
- acc: 0.7134 - val_loss: 0.7015 - val_acc: 0.6939

Epoch 7/30

7352/7352 [=====] - 95s 13ms/step - loss: 0.5692
- acc: 0.7477 - val_loss: 0.5995 - val_acc: 0.7387

Epoch 8/30

7352/7352 [=====] - 96s 13ms/step - loss: 0.4899
- acc: 0.7809 - val_loss: 0.5762 - val_acc: 0.7387

Epoch 9/30

7352/7352 [=====] - 90s 12ms/step - loss: 0.4482
- acc: 0.7886 - val_loss: 0.7413 - val_acc: 0.7126

Epoch 10/30

7352/7352 [=====] - 90s 12ms/step - loss: 0.4132
- acc: 0.8077 - val_loss: 0.5048 - val_acc: 0.7513

Epoch 11/30

7352/7352 [=====] - 89s 12ms/step - loss: 0.3985
- acc: 0.8274 - val_loss: 0.5234 - val_acc: 0.7452

Epoch 12/30

7352/7352 [=====] - 91s 12ms/step - loss: 0.3378
- acc: 0.8638 - val_loss: 0.4114 - val_acc: 0.8833

Epoch 13/30

7352/7352 [=====] - 91s 12ms/step - loss: 0.2947
- acc: 0.9051 - val_loss: 0.4386 - val_acc: 0.8731

Epoch 14/30

7352/7352 [=====] - 90s 12ms/step - loss: 0.2448
- acc: 0.9291 - val_loss: 0.3768 - val_acc: 0.8921

Epoch 15/30

7352/7352 [=====] - 91s 12ms/step - loss: 0.2157
- acc: 0.9331 - val_loss: 0.4441 - val_acc: 0.8931

Epoch 16/30

7352/7352 [=====] - 90s 12ms/step - loss: 0.2053
- acc: 0.9366 - val_loss: 0.4162 - val_acc: 0.8968

Epoch 17/30

7352/7352 [=====] - 89s 12ms/step - loss: 0.2028
- acc: 0.9404 - val_loss: 0.4538 - val_acc: 0.8962

Epoch 18/30

7352/7352 [=====] - 93s 13ms/step - loss: 0.1911
- acc: 0.9419 - val_loss: 0.3964 - val_acc: 0.8999

Epoch 19/30

7352/7352 [=====] - 96s 13ms/step - loss: 0.1912
- acc: 0.9407 - val_loss: 0.3165 - val_acc: 0.9030

Epoch 20/30

7352/7352 [=====] - 96s 13ms/step - loss: 0.1732
- acc: 0.9446 - val_loss: 0.4546 - val_acc: 0.8904

Epoch 21/30

7352/7352 [=====] - 94s 13ms/step - loss: 0.1782

- acc: 0.9444 - val_loss: 0.3346 - val_acc: 0.9063

Epoch 22/30

7352/7352 [=====] - 95s 13ms/step - loss: 0.1812

- acc: 0.9418 - val_loss: 0.8164 - val_acc: 0.8582

Epoch 23/30

7352/7352 [=====] - 95s 13ms/step - loss: 0.1824

- acc: 0.9426 - val_loss: 0.4240 - val_acc: 0.9036

Epoch 24/30

7352/7352 [=====] - 94s 13ms/step - loss: 0.1726

- acc: 0.9429 - val_loss: 0.4067 - val_acc: 0.9148

Epoch 25/30

7352/7352 [=====] - 96s 13ms/step - loss: 0.1737

- acc: 0.9411 - val_loss: 0.3396 - val_acc: 0.9074

Epoch 26/30

7352/7352 [=====] - 96s 13ms/step - loss: 0.1650

- acc: 0.9461 - val_loss: 0.3806 - val_acc: 0.9019

Epoch 27/30

7352/7352 [=====] - 89s 12ms/step - loss: 0.1925

- acc: 0.9415 - val_loss: 0.6464 - val_acc: 0.8850

Epoch 28/30

7352/7352 [=====] - 91s 12ms/step - loss: 0.1965

- acc: 0.9425 - val_loss: 0.3363 - val_acc: 0.9203

Epoch 29/30

7352/7352 [=====] - 92s 12ms/step - loss: 0.1889

- acc: 0.9431 - val_loss: 0.3737 - val_acc: 0.9158

Epoch 30/30

7352/7352 [=====] - 95s 13ms/step - loss: 0.1945

- acc: 0.9414 - val_loss: 0.3088 - val_acc: 0.9097

Out[0]:

<keras.callbacks.History at 0x29b5ee36a20>

In [0]:

```
# Confusion Matrix
print(confusion_matrix(Y_test, model.predict(X_test)))
```

Pred \ True	LAYING	SITTING	STANDING	WALKING	WALKING_DOWNSTAIRS	WALKING_UPSTAIRS
LAYING	512	0	25	0	0	0
SITTING	3	410	75	0	0	0
STANDING	0	87	445	0	0	0
WALKING	0	0	0	481	2	2
WALKING_DOWNSTAIRS	0	0	0	0	382	0
WALKING_UPSTAIRS	0	0	0	2	0	18

Pred \ True	WALKING_UPSTAIRS
LAYING	0
SITTING	3
STANDING	0
WALKING	13
WALKING_DOWNSTAIRS	38
WALKING_UPSTAIRS	451

In [0]:

```
score = model.evaluate(X_test, Y_test)
```

```
2947/2947 [=====] - 4s 2ms/step
```

In [0]:

```
score
```

Out[0]:

```
[0.3087582236972612, 0.9097387173396675]
```

- With a simple 2 layer architecture we got 90.09% accuracy and a loss of 0.30
- We can further improve the performance with Hyperparameter tuning

3.0 Hyperparameter tuning a single layered LSTM using KerasClassifier & Grid search

In [5]:

```
timesteps = len(X_train[0])
input_dim = len(X_train[0][0])
n_classes = _count_classes(Y_train)

print(timesteps)
print(input_dim)
print(len(X_train))
```

```
128
9
7352
```

In [0]:

```
# Credits: https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-python-keras/
# Function to create model, required for KerasClassifier
def create_model(cells=1,dropout_rate=0.0):
    # create model
    model = Sequential()
    model.add(LSTM(cells, input_shape=(timesteps, input_dim)))
    model.add(Dropout(dropout_rate))
    model.add(Dense(n_classes, activation='sigmoid'))
    # Compile model
    model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
    return model
```

In [0]:

```
model = KerasClassifier(build_fn=create_model, epochs=20, batch_size=50, verbose=0)
```

3.1 Grid Search

In [0]:

```
# defining the search parameters
import warnings
warnings.filterwarnings("ignore")
start = datetime.now()
cells=[64,128,150]
dropout_rate = [0.25, 0.35, 0.50]
param_grid = dict(cells= cells, dropout_rate=dropout_rate)
grid = GridSearchCV(estimator=model,param_grid=param_grid,cv=3)
grid_result = grid.fit(X_train, Y_train)
print('Time taken :', datetime.now() - start)
```

Time taken : 3:50:57.960481

3.2 Best estimator

In [0]:

```
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
```

Best: 0.653972 using {'cells': 64, 'dropout_rate': 0.35}

In [0]:

```
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

```
0.607726 (0.021996) with: {'cells': 64, 'dropout_rate': 0.25}
0.653972 (0.015658) with: {'cells': 64, 'dropout_rate': 0.35}
0.517818 (0.150945) with: {'cells': 64, 'dropout_rate': 0.5}
0.520947 (0.088254) with: {'cells': 128, 'dropout_rate': 0.25}
0.560800 (0.086814) with: {'cells': 128, 'dropout_rate': 0.35}
0.432127 (0.293107) with: {'cells': 128, 'dropout_rate': 0.5}
0.632345 (0.114650) with: {'cells': 150, 'dropout_rate': 0.25}
0.574129 (0.046813) with: {'cells': 150, 'dropout_rate': 0.35}
0.542029 (0.108520) with: {'cells': 150, 'dropout_rate': 0.5}
```

3.3 3-D Plot to visualize the metric for different values of hyperparameters

In [0]:

```
df=pd.DataFrame(grid.cv_results_)
df.head(2)
```

Out[0]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_cells	param_dropout_rate
0	473.090310	5.601252	5.374086	0.078879	64	0.25
1	476.258571	1.768062	5.696839	0.041143	64	0.35

In [0]:

```
df.to_csv('hyp.csv')
```

In [0]:

```
%matplotlib notebook
%matplotlib inline
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
import numpy as np
def enable_plotly_in_cell():
    import IPython
    from plotly.offline import init_notebook_mode
    display(IPython.core.display.HTML('''<script src="/static/components/requirejs/require.js"></script>'''))
    init_notebook_mode(connected=False)
```

In []:

```
# https://plot.ly/python/3d-axes/
#trace1 = go.Scatter3d(x=df['param_cells'],y=df['param_dropout_rate'],z=df['mean_test_score'], name = 'train')
trace2 = go.Scatter3d(x=df['param_cells'],y=df['param_dropout_rate'],z=df['mean_test_score'], name = 'Cross validation')
data = [trace2]
enable_plotly_in_cell()

layout = go.Layout(scene = dict(
    xaxis = dict(title='Number of LSTM cells'),
    yaxis = dict(title='Drop-out rate'),
    zaxis = dict(title='Accuracy'),))

fig = go.Figure(data=data, layout=layout)
offline.iplot(fig, filename='3d-scatter-colorscale')
```



3.4 Applying the best hyperparameters on the network

In [0]:

```
n_hidden= 64
dropout_rate= 0.35
```

Architecture

In [10]:

```
# Initiliazing the sequential model
model1 = Sequential()

model1.add(LSTM(n_hidden,input_shape=(timesteps, input_dim)))
model1.add(BatchNormalization())
model1.add(Dropout(dropout_rate))

model1.add(Dense(n_classes, activation='sigmoid'))
model1.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
lstm_2 (LSTM)	(None, 64)	18944
batch_normalization_2 (Batch Normalization)	(None, 64)	256
dropout_2 (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 6)	390
Total params: 19,590		
Trainable params: 19,462		
Non-trainable params: 128		

In [0]:

```
# Compiling the model
model1.compile(loss='categorical_crossentropy',optimizer='adam',metrics=['accuracy'])
```

3.5 Checkpointing the model and creating the callback list

In [0]:

```
from keras.callbacks import ModelCheckpoint
from keras.callbacks import CSVLogger
import matplotlib.pyplot as plt
from keras.callbacks import TensorBoard
import tensorflow as tf
import datetime
import keras

filepath="weights-{epoch:02d}-{val_accuracy:.2f}.hdf5"
checkpoints = ModelCheckpoint(filepath, monitor='val_accuracy', verbose=1, save_best_only=True, mode='max')
train_results = CSVLogger('train_results_2.log') #storing the training results in a pandas dataframe
callbacks_list = [checkpoints, train_results]
```

3.6 Fitting the model in batches

In [17]:

```
history= model1.fit(X_train,Y_train,batch_size=50,validation_data=(X_test, Y_test),nb_e  
poch=30,verbose=1,  
                    callbacks =callbacks_list)
```

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: UserWarning: The `nb_epoch` argument in `fit` has been renamed `epochs`.
```

Train on 7352 samples, validate on 2947 samples

Epoch 1/30

7352/7352 [=====] - 44s 6ms/step - loss: 0.8737 -
acc: 0.5903 - val_loss: 0.8794 - val_acc: 0.5148

Epoch 2/30

```
/usr/local/lib/python3.6/dist-packages/keras/callbacks.py:707: RuntimeWarning: Can save best model only with val_accuracy available, skipping.  
'skipping.' % (self.monitor), RuntimeWarning)
```

7352/7352 [=====] - 44s 6ms/step - loss: 0.7757 -
acc: 0.6091 - val_loss: 0.8388 - val_acc: 0.6257
Epoch 3/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.7354 -
acc: 0.6138 - val_loss: 0.8674 - val_acc: 0.5589
Epoch 4/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.7585 -
acc: 0.5871 - val_loss: 0.7672 - val_acc: 0.5809
Epoch 5/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.7604 -
acc: 0.5632 - val_loss: 0.8021 - val_acc: 0.5107
Epoch 6/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.7076 -
acc: 0.5717 - val_loss: 0.7230 - val_acc: 0.5701
Epoch 7/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.7145 -
acc: 0.5690 - val_loss: 0.7387 - val_acc: 0.5304
Epoch 8/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.7102 -
acc: 0.5690 - val_loss: 0.7256 - val_acc: 0.5073
Epoch 9/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.6796 -
acc: 0.5846 - val_loss: 0.7081 - val_acc: 0.6091
Epoch 10/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.7015 -
acc: 0.6204 - val_loss: 0.6679 - val_acc: 0.6637
Epoch 11/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.5928 -
acc: 0.6959 - val_loss: 0.6539 - val_acc: 0.6610
Epoch 12/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.6295 -
acc: 0.7047 - val_loss: 1.9109 - val_acc: 0.4917
Epoch 13/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.6305 -
acc: 0.7311 - val_loss: 0.5935 - val_acc: 0.7448
Epoch 14/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.4553 -
acc: 0.8402 - val_loss: 0.4621 - val_acc: 0.8626
Epoch 15/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.2337 -
acc: 0.9241 - val_loss: 0.6692 - val_acc: 0.8544
Epoch 16/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.1950 -
acc: 0.9300 - val_loss: 0.4736 - val_acc: 0.8758
Epoch 17/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.2293 -
acc: 0.9221 - val_loss: 0.5560 - val_acc: 0.8690
Epoch 18/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.2485 -
acc: 0.9144 - val_loss: 0.3738 - val_acc: 0.8907
Epoch 19/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.1847 -
acc: 0.9309 - val_loss: 0.2657 - val_acc: 0.9053
Epoch 20/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.1825 -
acc: 0.9338 - val_loss: 0.2923 - val_acc: 0.9128
Epoch 21/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.1516 -
acc: 0.9411 - val_loss: 0.2962 - val_acc: 0.9111
Epoch 22/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.1465 -

```

acc: 0.9436 - val_loss: 0.2487 - val_acc: 0.9074
Epoch 23/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.1433 -
acc: 0.9396 - val_loss: 0.3190 - val_acc: 0.9094
Epoch 24/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.1792 -
acc: 0.9310 - val_loss: 0.2996 - val_acc: 0.9121
Epoch 25/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.1683 -
acc: 0.9353 - val_loss: 0.3410 - val_acc: 0.8819
Epoch 26/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.1682 -
acc: 0.9377 - val_loss: 0.2552 - val_acc: 0.9019
Epoch 27/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.1430 -
acc: 0.9402 - val_loss: 0.2351 - val_acc: 0.9141
Epoch 28/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.1309 -
acc: 0.9444 - val_loss: 0.2480 - val_acc: 0.9040
Epoch 29/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.1266 -
acc: 0.9463 - val_loss: 0.2544 - val_acc: 0.9070
Epoch 30/30
7352/7352 [=====] - 43s 6ms/step - loss: 0.1200 -
acc: 0.9517 - val_loss: 0.2620 - val_acc: 0.9128

```

3.7 Confusion matrix

In [102]:

```

cm=confusion_matrix(Y_test, model1.predict(X_test))
cm

```

Out[102]:

	Pred	LAYING	SITTING	STANDING	WALKING	WALKING_DOWN
True						
LAYING		537	0	0	0	0
SITTING		0	371	116	1	0
STANDING		0	77	452	2	0
WALKING		0	8	0	467	15
WALKING_DOWNSTAIRS		0	0	0	4	412
WALKING_UPSTAIRS		0	0	0	13	7

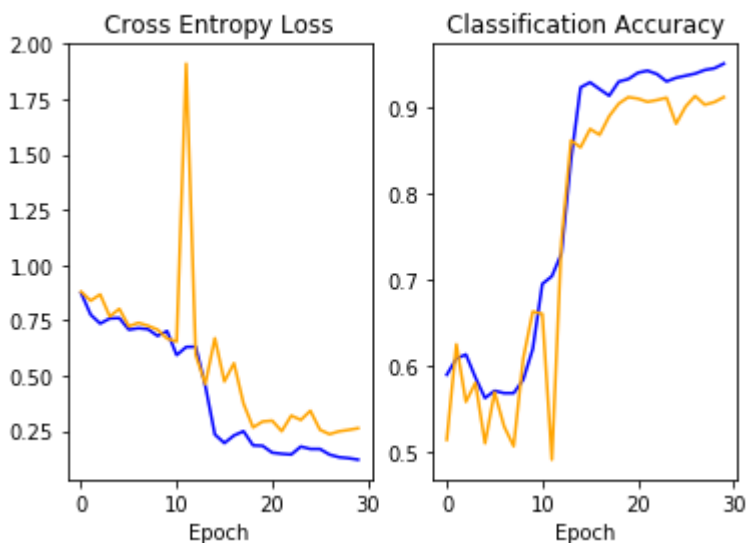
3.8 Plots on training results

In [0]:

```
# function to plot epoch vs Loss
%matplotlib notebook
%matplotlib inline
from matplotlib import pyplot
def plot(history):
    # plot loss
    pyplot.subplot(121)
    pyplot.title('Cross Entropy Loss')
    pyplot.xlabel('Epoch')
    pyplot.plot(history.history['loss'], color='blue', label='train')
    pyplot.plot(history.history['val_loss'], color='orange', label='test')
    # plot accuracy
    pyplot.subplot(122)
    pyplot.title('\nClassification Accuracy')
    pyplot.xlabel('Epoch')
    pyplot.plot(history.history['acc'], color='blue', label='train')
    pyplot.plot(history.history['val_acc'], color='orange', label='test')
```

In [41]:

```
plot(history)
```



3.9 Model Testing

In [20]:

```
score = model1.evaluate(X_test, Y_test, verbose=1)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

```
2947/2947 [=====] - 10s 3ms/step
Test loss: 0.2620189085359711
Test accuracy: 0.9127926705123854
```

4.0 Deep LSTM model

In [0]:

```
epochs = 50
batch_size= 50
n_hidden1 = 64
n_hidden2 =128
d1 = 0.50
d2 = 0.60 #using higher dropout rates
```

In [0]:

```
import keras.backend as K
K.clear_session()
```

4.1 Architecture

In [126]:

```
# Initiliazing the sequential model
model2 = Sequential()

model2.add(LSTM(n_hidden1,return_sequences=True,input_shape=(timesteps, input_dim)))
model2.add(BatchNormalization())
model2.add(Dropout(d1))

model2.add(LSTM(n_hidden2))
model2.add(BatchNormalization())
model2.add(Dropout(d2))

model2.add(Dense(n_classes, activation='sigmoid'))
model2.summary()
```

WARNING:tensorflow:Large dropout rate: 0.6 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
lstm_1 (LSTM)	(None, 128, 64)	18944
batch_normalization_1 (Batch Normalization)	(None, 128, 64)	256
dropout_1 (Dropout)	(None, 128, 64)	0
lstm_2 (LSTM)	(None, 128)	98816
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 6)	774
=====		
Total params: 119,302		
Trainable params: 118,918		
Non-trainable params: 384		

4.2 Compiling

In [0]:

```
# Compiling the model
model2.compile(loss='categorical_crossentropy',optimizer='adam',metrics=['accuracy'])
```

4.3 Checkpointing the model and creating the callback list

In [0]:

```
from keras.callbacks import ModelCheckpoint
from keras.callbacks import CSVLogger
import matplotlib.pyplot as plt
from keras.callbacks import TensorBoard
import tensorflow as tf
import datetime
import keras

filepath='model-ep{epoch:03d}-val_acc{val_acc:.3f}.h5'
checkpoints = ModelCheckpoint(filepath, monitor='val_accuracy', verbose=1, save_best_only=True, mode='max')
train_results = CSVLogger('train_results_model2.log') #storing the training results in a pandas dataframe
callbacks_list = [checkpoints, train_results]
```

4.4 Fitting the model in batches

In [129]:

```
# Fitting the model  
history1= model2.fit(X_train,Y_train,batch_size=batch_size,validation_data=(X_test, Y_t  
est),epochs=epochs)
```


Train on 7352 samples, validate on 2947 samples

Epoch 1/50

7352/7352 [=====] - 88s 12ms/step - loss: 1.0211
- acc: 0.6208 - val_loss: 0.8160 - val_acc: 0.6953

Epoch 2/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.7524
- acc: 0.6862 - val_loss: 0.7849 - val_acc: 0.6661

Epoch 3/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.7182
- acc: 0.6865 - val_loss: 0.7363 - val_acc: 0.7316

Epoch 4/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.6304
- acc: 0.7300 - val_loss: 0.9483 - val_acc: 0.6956

Epoch 5/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.5353
- acc: 0.8070 - val_loss: 0.5818 - val_acc: 0.8368

Epoch 6/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.3677
- acc: 0.8641 - val_loss: 0.4695 - val_acc: 0.8341

Epoch 7/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.2595
- acc: 0.8483 - val_loss: 0.4434 - val_acc: 0.7842

Epoch 8/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.2833
- acc: 0.8303 - val_loss: 0.3670 - val_acc: 0.7920

Epoch 9/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.2319
- acc: 0.8347 - val_loss: 0.4086 - val_acc: 0.7764

Epoch 10/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.2277
- acc: 0.8312 - val_loss: 0.3435 - val_acc: 0.8035

Epoch 11/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.2197
- acc: 0.8402 - val_loss: 0.3575 - val_acc: 0.7978

Epoch 12/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.2174
- acc: 0.8860 - val_loss: 0.3930 - val_acc: 0.9114

Epoch 13/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.1803
- acc: 0.9338 - val_loss: 0.4490 - val_acc: 0.8894

Epoch 14/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.1561
- acc: 0.9410 - val_loss: 0.4746 - val_acc: 0.8548

Epoch 15/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.1666
- acc: 0.9391 - val_loss: 0.2934 - val_acc: 0.9104

Epoch 16/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.1702
- acc: 0.9355 - val_loss: 0.3931 - val_acc: 0.8873

Epoch 17/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.1906
- acc: 0.9290 - val_loss: 0.3184 - val_acc: 0.9077

Epoch 18/50

7352/7352 [=====] - 87s 12ms/step - loss: 0.1631
- acc: 0.9361 - val_loss: 0.2773 - val_acc: 0.9135

Epoch 19/50

7352/7352 [=====] - 86s 12ms/step - loss: 0.1359
- acc: 0.9407 - val_loss: 0.3139 - val_acc: 0.9060

Epoch 20/50

7352/7352 [=====] - 86s 12ms/step - loss: 0.1375
- acc: 0.9479 - val_loss: 0.3403 - val_acc: 0.9097

Epoch 21/50
7352/7352 [=====] - 86s 12ms/step - loss: 0.1440
- acc: 0.9430 - val_loss: 0.3217 - val_acc: 0.9148

Epoch 22/50
7352/7352 [=====] - 85s 12ms/step - loss: 0.1313
- acc: 0.9484 - val_loss: 0.3400 - val_acc: 0.9097

Epoch 23/50
7352/7352 [=====] - 85s 12ms/step - loss: 0.1913
- acc: 0.9340 - val_loss: 0.2570 - val_acc: 0.9186

Epoch 24/50
7352/7352 [=====] - 86s 12ms/step - loss: 0.1379
- acc: 0.9412 - val_loss: 0.2645 - val_acc: 0.9281

Epoch 25/50
7352/7352 [=====] - 85s 12ms/step - loss: 0.1649
- acc: 0.9415 - val_loss: 0.2581 - val_acc: 0.9046

Epoch 26/50
7352/7352 [=====] - 85s 12ms/step - loss: 0.1326
- acc: 0.9478 - val_loss: 0.2355 - val_acc: 0.9355

Epoch 27/50
7352/7352 [=====] - 86s 12ms/step - loss: 0.1320
- acc: 0.9490 - val_loss: 0.2499 - val_acc: 0.9253

Epoch 28/50
7352/7352 [=====] - 87s 12ms/step - loss: 0.1220
- acc: 0.9489 - val_loss: 0.2754 - val_acc: 0.9257

Epoch 29/50
7352/7352 [=====] - 87s 12ms/step - loss: 0.1227
- acc: 0.9486 - val_loss: 0.2694 - val_acc: 0.9209

Epoch 30/50
7352/7352 [=====] - 88s 12ms/step - loss: 0.1287
- acc: 0.9463 - val_loss: 0.2407 - val_acc: 0.9281

Epoch 31/50
7352/7352 [=====] - 88s 12ms/step - loss: 0.1671
- acc: 0.9306 - val_loss: 0.2330 - val_acc: 0.9175

Epoch 32/50
7352/7352 [=====] - 88s 12ms/step - loss: 0.1439
- acc: 0.9369 - val_loss: 0.3069 - val_acc: 0.9074

Epoch 33/50
7352/7352 [=====] - 88s 12ms/step - loss: 0.1429
- acc: 0.9392 - val_loss: 0.3173 - val_acc: 0.9104

Epoch 34/50
7352/7352 [=====] - 88s 12ms/step - loss: 0.1222
- acc: 0.9506 - val_loss: 0.2809 - val_acc: 0.9318

Epoch 35/50
7352/7352 [=====] - 88s 12ms/step - loss: 0.1227
- acc: 0.9521 - val_loss: 0.2797 - val_acc: 0.9233

Epoch 36/50
7352/7352 [=====] - 88s 12ms/step - loss: 0.1186
- acc: 0.9504 - val_loss: 0.3137 - val_acc: 0.9226

Epoch 37/50
7352/7352 [=====] - 88s 12ms/step - loss: 0.1596
- acc: 0.9354 - val_loss: 0.3006 - val_acc: 0.9128

Epoch 38/50
7352/7352 [=====] - 88s 12ms/step - loss: 0.1533
- acc: 0.9351 - val_loss: 0.3289 - val_acc: 0.8965

Epoch 39/50
7352/7352 [=====] - 88s 12ms/step - loss: 0.1540
- acc: 0.9370 - val_loss: 0.2790 - val_acc: 0.9243

Epoch 40/50
7352/7352 [=====] - 87s 12ms/step - loss: 0.1287
- acc: 0.9464 - val_loss: 0.2605 - val_acc: 0.9284

Epoch 41/50

```

7352/7352 [=====] - 88s 12ms/step - loss: 0.1227
- acc: 0.9479 - val_loss: 0.2856 - val_acc: 0.9260
Epoch 42/50
7352/7352 [=====] - 87s 12ms/step - loss: 0.1214
- acc: 0.9467 - val_loss: 0.3178 - val_acc: 0.9274
Epoch 43/50
7352/7352 [=====] - 87s 12ms/step - loss: 0.1218
- acc: 0.9493 - val_loss: 0.3100 - val_acc: 0.9270
Epoch 44/50
7352/7352 [=====] - 87s 12ms/step - loss: 0.1222
- acc: 0.9497 - val_loss: 0.3382 - val_acc: 0.9182
Epoch 45/50
7352/7352 [=====] - 89s 12ms/step - loss: 0.1255
- acc: 0.9509 - val_loss: 0.3199 - val_acc: 0.9230
Epoch 46/50
7352/7352 [=====] - 89s 12ms/step - loss: 0.1120
- acc: 0.9532 - val_loss: 0.3275 - val_acc: 0.9213
Epoch 47/50
7352/7352 [=====] - 87s 12ms/step - loss: 0.1225
- acc: 0.9487 - val_loss: 0.3052 - val_acc: 0.9247
Epoch 48/50
7352/7352 [=====] - 88s 12ms/step - loss: 0.1304
- acc: 0.9421 - val_loss: 0.3078 - val_acc: 0.9165
Epoch 49/50
7352/7352 [=====] - 87s 12ms/step - loss: 0.1237
- acc: 0.9484 - val_loss: 0.3364 - val_acc: 0.9186
Epoch 50/50
7352/7352 [=====] - 87s 12ms/step - loss: 0.1196
- acc: 0.9524 - val_loss: 0.3126 - val_acc: 0.9308

```

4.5 Confusion matrix

In [132]:

```

cm1= confusion_matrix(Y_test, model2.predict(X_test))
cm1

```


Out[132]:

	Pred	LAYING	SITTING	STANDING	WALKING	WALKING_DOWN
True						
LAYING		537	0	0	0	0
SITTING		0	370	118	0	0
STANDING		0	50	482	0	0
WALKING		0	0	0	466	27
WALKING_DOWNSTAIRS		0	0	0	1	418
WALKING_UPSTAIRS		0	0	0	1	0

4.6 Plots on training results

In [134]:

```
plot(history1)
```



4.7 Model Testing

In [135]:

```
score = model2.evaluate(X_test, Y_test, verbose=1)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

```
2947/2947 [=====] - 19s 6ms/step
Test loss: 0.31255650963575987
Test accuracy: 0.9307770614183916
```

5.0 Summary

In [3]:

```
#Ref: http://zetcode.com/python/prettytable/
```

```
from prettytable import PrettyTable
x=PrettyTable()
x.field_names=["Model","Test loss","Test accuracy"]

x.add_row(["1 layered LSTM without hyp tuning","0.3088","90.97%"])
x.add_row(["1 layered LSTM with hyp tuning","0.2620","91.30%"])
x.add_row(["Deep 2 layered LSTM","0.3126","93.08%"])

print(x)
```

Model	Test loss	Test accuracy
1 layered LSTM without hyp tuning	0.3088	90.97%
1 layered LSTM with hyp tuning	0.2620	91.30%
Deep 2 layered LSTM	0.3126	93.08%