

MLP's on MNIST using Keras

1.0 Loading the MNIST data

In [1]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use th
is command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

In [2]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

2.0 Splitting the data

In [3]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 0s 0us/step

In [4]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%
d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d
, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

In [5]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [6]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

```
Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)
```

In [7]:

```
# An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0 18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
```

In [8]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 
```

```
X_train = X_train/255
X_test = X_test/255
```

In [9]:

```
# example data point after normlizing  
print(X_train[0])
```

[illegible]

0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.21568627	0.6745098	0.88627451	0.99215686	0.99215686	0.99215686
0.99215686	0.95686275	0.52156863	0.04313725	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.53333333	0.99215686
0.99215686	0.99215686	0.83137255	0.52941176	0.51764706	0.0627451
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

In [10]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

Class label of first image : 5

After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

3.0 Softmax classifier

In [0]:

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT.X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation
```


In [0]:

```
# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

In [0]:

```
# start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
```

In [0]:

```
# Before training a model, you need to configure the learning process, which is done via the compile method

# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer , https://keras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize., https://keras.io/losses/
# A list of metrics. For any classification problem you will want to set this to metric_s=['accuracy']. https://keras.io/metrics/

# Note: when using the categorical_crossentropy loss, your targets should be in categorical format
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that is all-zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted our labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None, validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, steps_per_epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss values and
# metrics values at successive epochs, as well as validation loss values and validation metrics values (if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 49us/step - loss: 1.2935
- acc: 0.6829 - val_loss: 0.8171 - val_acc: 0.8312

Epoch 2/20

60000/60000 [=====] - 2s 35us/step - loss: 0.7213
- acc: 0.8385 - val_loss: 0.6099 - val_acc: 0.8623

Epoch 3/20

60000/60000 [=====] - 2s 35us/step - loss: 0.5904
- acc: 0.8584 - val_loss: 0.5275 - val_acc: 0.8741

Epoch 4/20

60000/60000 [=====] - 2s 35us/step - loss: 0.5278
- acc: 0.8682 - val_loss: 0.4815 - val_acc: 0.8814

Epoch 5/20

60000/60000 [=====] - 2s 35us/step - loss: 0.4897
- acc: 0.8747 - val_loss: 0.4515 - val_acc: 0.8870

Epoch 6/20

34432/60000 [=====>.....] - ETA: 0s - loss: 0.4697 - acc: 0.8773
60000/60000 [=====] - 2s 35us/step - loss: 0.4635 - acc: 0.8799 - val_loss: 0.4303 - val_acc: 0.8898

Epoch 7/20

60000/60000 [=====] - 2s 35us/step - loss: 0.4442
- acc: 0.8837 - val_loss: 0.4138 - val_acc: 0.8917

Epoch 8/20

60000/60000 [=====] - 2s 35us/step - loss: 0.4290
- acc: 0.8866 - val_loss: 0.4010 - val_acc: 0.8952

Epoch 9/20

60000/60000 [=====] - 2s 35us/step - loss: 0.4169
- acc: 0.8894 - val_loss: 0.3909 - val_acc: 0.8971

Epoch 10/20

60000/60000 [=====] - 2s 35us/step - loss: 0.4067
- acc: 0.8911 - val_loss: 0.3818 - val_acc: 0.8994

Epoch 11/20

60000/60000 [=====] - 2s 35us/step - loss: 0.3982
- acc: 0.8926 - val_loss: 0.3743 - val_acc: 0.9006

Epoch 12/20

1792/60000 [.....] - ETA: 1s - loss: 0.4077 - acc: 0.8895
60000/60000 [=====] - 2s 35us/step - loss: 0.3908 - acc: 0.8948 - val_loss: 0.3681 - val_acc: 0.9020

Epoch 13/20

60000/60000 [=====] - 2s 35us/step - loss: 0.3844
- acc: 0.8960 - val_loss: 0.3624 - val_acc: 0.9039

Epoch 14/20

60000/60000 [=====] - 2s 35us/step - loss: 0.3786
- acc: 0.8972 - val_loss: 0.3575 - val_acc: 0.9046

Epoch 15/20

60000/60000 [=====] - 2s 35us/step - loss: 0.3735
- acc: 0.8981 - val_loss: 0.3528 - val_acc: 0.9058

Epoch 16/20

60000/60000 [=====] - 2s 35us/step - loss: 0.3689
- acc: 0.8993 - val_loss: 0.3490 - val_acc: 0.9062

Epoch 17/20

60000/60000 [=====] - 2s 35us/step - loss: 0.3648
- acc: 0.9004 - val_loss: 0.3455 - val_acc: 0.9066

Epoch 18/20

60000/60000 [=====] - 2s 35us/step - loss: 0.3610
- acc: 0.9016 - val_loss: 0.3419 - val_acc: 0.9077

Epoch 19/20

60000/60000 [=====] - 2s 35us/step - loss: 0.3575
- acc: 0.9024 - val_loss: 0.3389 - val_acc: 0.9088

Epoch 20/20

```
60000/60000 [=====] - 2s 35us/step - loss: 0.3544
- acc: 0.9032 - val_loss: 0.3362 - val_acc: 0.9093
```

In [0]:

```
score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.3362289469957352

Test accuracy: 0.9093

4.0 MLP on 784-512-128-10 architecture

4.1 MLP + Sigmoid activation + SGD Optimizer

In [0]:

```
# Multilayer perceptron

model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 512)	401920
dense_3 (Dense)	(None, 128)	65664
dense_4 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

In [0]:

```
model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 42us/step - loss: 2.2708
- acc: 0.2169 - val_loss: 2.2197 - val_acc: 0.2768

Epoch 2/20

60000/60000 [=====] - 2s 42us/step - loss: 2.1739
- acc: 0.4237 - val_loss: 2.1143 - val_acc: 0.4877

Epoch 3/20

60000/60000 [=====] - 2s 42us/step - loss: 2.0506
- acc: 0.5485 - val_loss: 1.9659 - val_acc: 0.5524

Epoch 4/20

60000/60000 [=====] - 3s 42us/step - loss: 1.8796
- acc: 0.6135 - val_loss: 1.7673 - val_acc: 0.6481

Epoch 5/20

60000/60000 [=====] - 3s 42us/step - loss: 1.6674
- acc: 0.6659 - val_loss: 1.5414 - val_acc: 0.7205

Epoch 6/20

5376/60000 [=>.....] - ETA: 2s - loss: 1.5430 - acc: 0.698160000/60000 [=====] - 2s 41us/step - loss: 1.4449 - acc: 0.7125 - val_loss: 1.3233 - val_acc: 0.7513

Epoch 7/20

60000/60000 [=====] - 2s 41us/step - loss: 1.2442
- acc: 0.7466 - val_loss: 1.1406 - val_acc: 0.7599

Epoch 8/20

60000/60000 [=====] - 2s 41us/step - loss: 1.0815
- acc: 0.7722 - val_loss: 0.9974 - val_acc: 0.7968

Epoch 9/20

60000/60000 [=====] - 2s 41us/step - loss: 0.9544
- acc: 0.7940 - val_loss: 0.8867 - val_acc: 0.8060

Epoch 10/20

60000/60000 [=====] - 2s 41us/step - loss: 0.8556
- acc: 0.8082 - val_loss: 0.7984 - val_acc: 0.8227

Epoch 11/20

31232/60000 [=====>.....] - ETA: 1s - loss: 0.7920 - acc: 0.820760000/60000 [=====] - 2s 42us/step - loss: 0.7776 - acc: 0.8225 - val_loss: 0.7292 - val_acc: 0.8335

Epoch 12/20

60000/60000 [=====] - 2s 41us/step - loss: 0.7154
- acc: 0.8333 - val_loss: 0.6741 - val_acc: 0.8422

Epoch 13/20

60000/60000 [=====] - 2s 41us/step - loss: 0.6650
- acc: 0.8412 - val_loss: 0.6282 - val_acc: 0.8500

Epoch 14/20

60000/60000 [=====] - 2s 41us/step - loss: 0.6237
- acc: 0.8485 - val_loss: 0.5906 - val_acc: 0.8585

Epoch 15/20

60000/60000 [=====] - 2s 41us/step - loss: 0.5893
- acc: 0.8546 - val_loss: 0.5591 - val_acc: 0.8616

Epoch 16/20

34432/60000 [=====>.....] - ETA: 0s - loss: 0.5691 - acc: 0.857860000/60000 [=====] - 2s 41us/step - loss: 0.5606 - acc: 0.8599 - val_loss: 0.5329 - val_acc: 0.8672

Epoch 17/20

60000/60000 [=====] - 2s 41us/step - loss: 0.5361
- acc: 0.8641 - val_loss: 0.5095 - val_acc: 0.8703

Epoch 18/20

60000/60000 [=====] - 2s 42us/step - loss: 0.5151
- acc: 0.8676 - val_loss: 0.4904 - val_acc: 0.8736

Epoch 19/20

60000/60000 [=====] - 2s 41us/step - loss: 0.4969
- acc: 0.8710 - val_loss: 0.4732 - val_acc: 0.8782

Epoch 20/20

60000/60000 [=====] - 2s 42us/step - loss: 0.4810

- acc: 0.8739 - val_loss: 0.4583 - val_acc: 0.8816

In [0]:

```
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.4582893396139145

Test accuracy: 0.8816

4.2 MLP + Sigmoid activation + ADAM

In [0]:

```
model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 512)	401920
dense_6 (Dense)	(None, 128)	65664
dense_7 (Dense)	(None, 10)	1290

=====

Total params: 468,874

Trainable params: 468,874

Non-trainable params: 0

=====

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 55us/step - loss: 0.5308
- acc: 0.8636 - val_loss: 0.2550 - val_acc: 0.9259

Epoch 2/20

53120/60000 [=====>....] - ETA: 0s - loss: 0.2244 - acc: 0.9340
60000/60000 [=====] - 3s 51us/step - loss: 0.2205 - acc: 0.9351 - val_loss: 0.1946 - val_acc: 0.9417

Epoch 3/20

60000/60000 [=====] - 3s 51us/step - loss: 0.1650
- acc: 0.9512 - val_loss: 0.1421 - val_acc: 0.9570

Epoch 4/20

60000/60000 [=====] - 3s 51us/step - loss: 0.1279
- acc: 0.9614 - val_loss: 0.1238 - val_acc: 0.9645

Epoch 5/20

60000/60000 [=====] - 3s 51us/step - loss: 0.1010
- acc: 0.9704 - val_loss: 0.1029 - val_acc: 0.9693

Epoch 6/20

60000/60000 [=====] - 3s 51us/step - loss: 0.0801
- acc: 0.9763 - val_loss: 0.0877 - val_acc: 0.9725

Epoch 7/20

4480/60000 [=>.....] - ETA: 2s - loss: 0.0632 - acc: 0.9795
60000/60000 [=====] - 3s 51us/step - loss: 0.0645 - acc: 0.9809 - val_loss: 0.0831 - val_acc: 0.9751

Epoch 8/20

60000/60000 [=====] - 3s 51us/step - loss: 0.0519
- acc: 0.9842 - val_loss: 0.0724 - val_acc: 0.9780

Epoch 9/20

60000/60000 [=====] - 3s 51us/step - loss: 0.0433
- acc: 0.9872 - val_loss: 0.0714 - val_acc: 0.9786

Epoch 10/20

60000/60000 [=====] - 3s 51us/step - loss: 0.0347
- acc: 0.9898 - val_loss: 0.0695 - val_acc: 0.9776

Epoch 11/20

60000/60000 [=====] - 3s 51us/step - loss: 0.0268
- acc: 0.9930 - val_loss: 0.0659 - val_acc: 0.9796

Epoch 12/20

60000/60000 [=====] - 3s 52us/step - loss: 0.0219
- acc: 0.9944 - val_loss: 0.0642 - val_acc: 0.9809

Epoch 13/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0180
- acc: 0.9953 - val_loss: 0.0677 - val_acc: 0.9794

Epoch 14/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0133
- acc: 0.9970 - val_loss: 0.0647 - val_acc: 0.9803

Epoch 15/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0114
- acc: 0.9975 - val_loss: 0.0628 - val_acc: 0.9812

Epoch 16/20

57344/60000 [=====>..] - ETA: 0s - loss: 0.0085 - acc: 0.998260000/60000 [=====] - 3s 50us/step - loss: 0.0085 - acc: 0.9982 - val_loss: 0.0666 - val_acc: 0.9806

Epoch 17/20

60000/60000 [=====] - 3s 51us/step - loss: 0.0070 - acc: 0.9986 - val_loss: 0.0643 - val_acc: 0.9822

Epoch 18/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0061 - acc: 0.9986 - val_loss: 0.0656 - val_acc: 0.9818

Epoch 19/20

60000/60000 [=====] - 3s 51us/step - loss: 0.0055 - acc: 0.9988 - val_loss: 0.0811 - val_acc: 0.9774

Epoch 20/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0038 - acc: 0.9992 - val_loss: 0.0723 - val_acc: 0.9818

In [0]:

```
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06385514608082886

Test accuracy: 0.9824



4.3 MLP + ReLU +SGD

In [0]:

```
# Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, \sigma) = N(0, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.125)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
dense_9 (Dense)	(None, 128)	65664
dense_10 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

In [0]:

```
model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])  
  
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 67us/step - loss: 0.7579
- acc: 0.7812 - val_loss: 0.3951 - val_acc: 0.8921

Epoch 2/20

60000/60000 [=====] - 4s 64us/step - loss: 0.3535
- acc: 0.8998 - val_loss: 0.3040 - val_acc: 0.9153

Epoch 3/20

60000/60000 [=====] - 4s 64us/step - loss: 0.2900
- acc: 0.9172 - val_loss: 0.2648 - val_acc: 0.9253

Epoch 4/20

60000/60000 [=====] - 4s 60us/step - loss: 0.2558
- acc: 0.9269 - val_loss: 0.2393 - val_acc: 0.9316

Epoch 5/20

60000/60000 [=====] - 4s 58us/step - loss: 0.2324
- acc: 0.9340 - val_loss: 0.2210 - val_acc: 0.9371

Epoch 6/20

60000/60000 [=====] - 4s 64us/step - loss: 0.2144
- acc: 0.9391 - val_loss: 0.2072 - val_acc: 0.9400

Epoch 7/20

60000/60000 [=====] - 4s 66us/step - loss: 0.1995
- acc: 0.9443 - val_loss: 0.1957 - val_acc: 0.9444

Epoch 8/20

60000/60000 [=====] - 4s 61us/step - loss: 0.1872
- acc: 0.9476 - val_loss: 0.1848 - val_acc: 0.9456

Epoch 9/20

60000/60000 [=====] - 3s 57us/step - loss: 0.1763
- acc: 0.9507 - val_loss: 0.1771 - val_acc: 0.9488

Epoch 10/20

60000/60000 [=====] - 4s 60us/step - loss: 0.1668
- acc: 0.9539 - val_loss: 0.1682 - val_acc: 0.9506

Epoch 11/20

60000/60000 [=====] - 4s 71us/step - loss: 0.1585
- acc: 0.9560 - val_loss: 0.1623 - val_acc: 0.9518

Epoch 12/20

60000/60000 [=====] - 7s 113us/step - loss: 0.151
1 - acc: 0.9577 - val_loss: 0.1560 - val_acc: 0.9543

Epoch 13/20

60000/60000 [=====] - 7s 115us/step - loss: 0.144
3 - acc: 0.9596 - val_loss: 0.1517 - val_acc: 0.9557

Epoch 14/20

60000/60000 [=====] - 7s 111us/step - loss: 0.137
9 - acc: 0.9615 - val_loss: 0.1474 - val_acc: 0.9572

Epoch 15/20

60000/60000 [=====] - 4s 66us/step - loss: 0.1323
- acc: 0.9628 - val_loss: 0.1429 - val_acc: 0.9580

Epoch 16/20

60000/60000 [=====] - 4s 69us/step - loss: 0.1270
- acc: 0.9645 - val_loss: 0.1371 - val_acc: 0.9598

Epoch 17/20

60000/60000 [=====] - 7s 110us/step - loss: 0.122
1 - acc: 0.9661 - val_loss: 0.1351 - val_acc: 0.9602

Epoch 18/20

60000/60000 [=====] - 4s 62us/step - loss: 0.1177
- acc: 0.9671 - val_loss: 0.1309 - val_acc: 0.9618

Epoch 19/20

60000/60000 [=====] - 4s 60us/step - loss: 0.1136
- acc: 0.9685 - val_loss: 0.1263 - val_acc: 0.9631

Epoch 20/20

60000/60000 [=====] - 5s 79us/step - loss: 0.1094
- acc: 0.9694 - val_loss: 0.1241 - val_acc: 0.9631

In [0]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.12405014228336513

Test accuracy: 0.9631



In [0]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



4.4 MLP + ReLU + ADAM

In [0]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 512)	401920
dense_12 (Dense)	(None, 128)	65664
dense_13 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 121us/step - loss: 0.234

1 - acc: 0.9295 - val_loss: 0.1165 - val_acc: 0.9652

Epoch 2/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0878

- acc: 0.9729 - val_loss: 0.0883 - val_acc: 0.9720

Epoch 3/20

60000/60000 [=====] - 5s 75us/step - loss: 0.0544

- acc: 0.9825 - val_loss: 0.0860 - val_acc: 0.9729

Epoch 4/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0354

- acc: 0.9885 - val_loss: 0.0699 - val_acc: 0.9797

Epoch 5/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0266

- acc: 0.9914 - val_loss: 0.0720 - val_acc: 0.9788

Epoch 6/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0200

- acc: 0.9941 - val_loss: 0.0696 - val_acc: 0.9803

Epoch 7/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0155

- acc: 0.9951 - val_loss: 0.0640 - val_acc: 0.9829

Epoch 8/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0140

- acc: 0.9952 - val_loss: 0.0848 - val_acc: 0.9792

Epoch 9/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0143

- acc: 0.9952 - val_loss: 0.0837 - val_acc: 0.9796

Epoch 10/20

60000/60000 [=====] - 7s 115us/step - loss: 0.012

8 - acc: 0.9958 - val_loss: 0.0946 - val_acc: 0.9782

Epoch 11/20

60000/60000 [=====] - 7s 125us/step - loss: 0.008

1 - acc: 0.9974 - val_loss: 0.0682 - val_acc: 0.9826

Epoch 12/20

60000/60000 [=====] - 8s 129us/step - loss: 0.012

1 - acc: 0.9959 - val_loss: 0.0793 - val_acc: 0.9816

Epoch 13/20

60000/60000 [=====] - 8s 133us/step - loss: 0.010

7 - acc: 0.9963 - val_loss: 0.0746 - val_acc: 0.9820

Epoch 14/20

60000/60000 [=====] - 8s 129us/step - loss: 0.011

3 - acc: 0.9960 - val_loss: 0.0813 - val_acc: 0.9816

Epoch 15/20

60000/60000 [=====] - 5s 77us/step - loss: 0.0058

- acc: 0.9982 - val_loss: 0.0770 - val_acc: 0.9842

Epoch 16/20

```

60000/60000 [=====] - 4s 65us/step - loss: 0.0040
- acc: 0.9987 - val_loss: 0.0930 - val_acc: 0.9808
Epoch 17/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0119
- acc: 0.9959 - val_loss: 0.0813 - val_acc: 0.9819
Epoch 18/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0105
- acc: 0.9966 - val_loss: 0.1000 - val_acc: 0.9803
Epoch 19/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0064
- acc: 0.9981 - val_loss: 0.0852 - val_acc: 0.9831
Epoch 20/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0056
- acc: 0.9982 - val_loss: 0.1029 - val_acc: 0.9805

```

In [0]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')


# List of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.10294274219236926

Test accuracy: 0.9805



In [0]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



4.5 MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

In [0]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(\theta, \sigma) = N(\theta, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(\theta, \sigma) = N(\theta, 0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Layer (type)	Output Shape	Param #
dense_14 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dense_15 (Dense)	(None, 128)	65664
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dense_16 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

In [0]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```


Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 138us/step - loss: 0.303

6 - acc: 0.9104 - val_loss: 0.2116 - val_acc: 0.9376

Epoch 2/20

60000/60000 [=====] - 10s 170us/step - loss: 0.17

47 - acc: 0.9483 - val_loss: 0.1670 - val_acc: 0.9505

Epoch 3/20

60000/60000 [=====] - 13s 220us/step - loss: 0.13

67 - acc: 0.9599 - val_loss: 0.1451 - val_acc: 0.9567

Epoch 4/20

60000/60000 [=====] - 9s 156us/step - loss: 0.113

4 - acc: 0.9666 - val_loss: 0.1335 - val_acc: 0.9603

Epoch 5/20

60000/60000 [=====] - 13s 211us/step - loss: 0.09

49 - acc: 0.9703 - val_loss: 0.1325 - val_acc: 0.9589

Epoch 6/20

60000/60000 [=====] - 7s 119us/step - loss: 0.080

2 - acc: 0.9758 - val_loss: 0.1139 - val_acc: 0.9652

Epoch 7/20

60000/60000 [=====] - 8s 127us/step - loss: 0.068

2 - acc: 0.9787 - val_loss: 0.1136 - val_acc: 0.9666

Epoch 8/20

60000/60000 [=====] - 7s 124us/step - loss: 0.060

8 - acc: 0.9815 - val_loss: 0.1114 - val_acc: 0.9666

Epoch 9/20

60000/60000 [=====] - 8s 129us/step - loss: 0.053

2 - acc: 0.9837 - val_loss: 0.1167 - val_acc: 0.9666

Epoch 10/20

60000/60000 [=====] - 7s 123us/step - loss: 0.045

5 - acc: 0.9856 - val_loss: 0.0962 - val_acc: 0.9718

Epoch 11/20

60000/60000 [=====] - 7s 112us/step - loss: 0.037

6 - acc: 0.9880 - val_loss: 0.1102 - val_acc: 0.9673

Epoch 12/20

60000/60000 [=====] - 7s 124us/step - loss: 0.035

0 - acc: 0.9889 - val_loss: 0.1033 - val_acc: 0.9710

Epoch 13/20

60000/60000 [=====] - 7s 124us/step - loss: 0.030

8 - acc: 0.9903 - val_loss: 0.1020 - val_acc: 0.9712

Epoch 14/20

60000/60000 [=====] - 7s 123us/step - loss: 0.027

1 - acc: 0.9913 - val_loss: 0.1038 - val_acc: 0.9727

Epoch 15/20

60000/60000 [=====] - 7s 122us/step - loss: 0.023

1 - acc: 0.9926 - val_loss: 0.1019 - val_acc: 0.9717

Epoch 16/20

60000/60000 [=====] - 8s 127us/step - loss: 0.022

0 - acc: 0.9928 - val_loss: 0.1110 - val_acc: 0.9703

Epoch 17/20

60000/60000 [=====] - 7s 114us/step - loss: 0.022

9 - acc: 0.9928 - val_loss: 0.1067 - val_acc: 0.9739

Epoch 18/20

60000/60000 [=====] - 8s 128us/step - loss: 0.020

3 - acc: 0.9935 - val_loss: 0.0982 - val_acc: 0.9738

Epoch 19/20

60000/60000 [=====] - 7s 125us/step - loss: 0.017

1 - acc: 0.9944 - val_loss: 0.1056 - val_acc: 0.9706

Epoch 20/20

60000/60000 [=====] - 11s 182us/step - loss: 0.01

46 - acc: 0.9952 - val_loss: 0.1046 - val_acc: 0.9732

In [0]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10456635547156475

Test accuracy: 0.9732



In [0]:


```
w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



4.6 MLP + Dropout + AdamOptimizer

In [0]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_17 (Dense)	(None, 512)	401920
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_1 (Dropout)	(None, 512)	0
dense_18 (Dense)	(None, 128)	65664
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_19 (Dense)	(None, 10)	1290
=====	=====	=====
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		
=====	=====	=====

In [0]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 14s 227us/step - loss: 0.66

12 - acc: 0.7951 - val_loss: 0.2860 - val_acc: 0.9166

Epoch 2/20

60000/60000 [=====] - 8s 136us/step - loss: 0.425

0 - acc: 0.8710 - val_loss: 0.2545 - val_acc: 0.9252

Epoch 3/20

60000/60000 [=====] - 12s 198us/step - loss: 0.38

41 - acc: 0.8846 - val_loss: 0.2391 - val_acc: 0.9298

Epoch 4/20

60000/60000 [=====] - 8s 138us/step - loss: 0.355

1 - acc: 0.8927 - val_loss: 0.2279 - val_acc: 0.9325

Epoch 5/20

60000/60000 [=====] - 7s 123us/step - loss: 0.335

5 - acc: 0.8986 - val_loss: 0.2127 - val_acc: 0.9356

Epoch 6/20

60000/60000 [=====] - 8s 136us/step - loss: 0.323

4 - acc: 0.9031 - val_loss: 0.2029 - val_acc: 0.9387: 1s - loss:

Epoch 7/20

60000/60000 [=====] - 8s 131us/step - loss: 0.306

8 - acc: 0.9077 - val_loss: 0.1927 - val_acc: 0.9421

Epoch 8/20

60000/60000 [=====] - 11s 185us/step - loss: 0.29

33 - acc: 0.9113 - val_loss: 0.1836 - val_acc: 0.9453

Epoch 9/20

60000/60000 [=====] - 13s 222us/step - loss: 0.28

50 - acc: 0.9131 - val_loss: 0.1797 - val_acc: 0.9451

Epoch 10/20

60000/60000 [=====] - 14s 236us/step - loss: 0.27

15 - acc: 0.9187 - val_loss: 0.1738 - val_acc: 0.9465

Epoch 11/20

60000/60000 [=====] - 8s 141us/step - loss: 0.261

1 - acc: 0.9214 - val_loss: 0.1671 - val_acc: 0.9506

Epoch 12/20

60000/60000 [=====] - 8s 134us/step - loss: 0.246

4 - acc: 0.9252 - val_loss: 0.1554 - val_acc: 0.9525

Epoch 13/20

60000/60000 [=====] - 8s 137us/step - loss: 0.238

2 - acc: 0.9278 - val_loss: 0.1479 - val_acc: 0.9554

Epoch 14/20

60000/60000 [=====] - 8s 136us/step - loss: 0.227

5 - acc: 0.9313 - val_loss: 0.1375 - val_acc: 0.9580

Epoch 15/20

60000/60000 [=====] - 8s 137us/step - loss: 0.218

3 - acc: 0.9337 - val_loss: 0.1326 - val_acc: 0.9599

Epoch 16/20

60000/60000 [=====] - 8s 138us/step - loss: 0.206

8 - acc: 0.9384 - val_loss: 0.1297 - val_acc: 0.9613 loss: 0.2066 - ac

Epoch 17/20

60000/60000 [=====] - 8s 139us/step - loss: 0.201

1 - acc: 0.9395 - val_loss: 0.1181 - val_acc: 0.9646

Epoch 18/20

60000/60000 [=====] - 8s 137us/step - loss: 0.188

6 - acc: 0.9435 - val_loss: 0.1145 - val_acc: 0.9658

Epoch 19/20

60000/60000 [=====] - 8s 138us/step - loss: 0.182

1 - acc: 0.9451 - val_loss: 0.1104 - val_acc: 0.9662

Epoch 20/20

60000/60000 [=====] - 8s 139us/step - loss: 0.173

9 - acc: 0.9473 - val_loss: 0.1093 - val_acc: 0.9679

In [0]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig, ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1093290721397847

Test accuracy: 0.9679



In [0]:


```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



4.7 Hyper-parameter tuning of Keras models using Sklearn

In [0]:

```

from keras.optimizers import Adam,RMSprop,SGD
def best_hyperparameters(activ):

    model = Sequential()
    model.add(Dense(512, activation=activ, input_shape=(input_dim,), kernel_initializer
=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(Dense(128, activation=activ, kernel_initializer=RandomNormal(mean=0.0, st
ddev=0.125, seed=None)) )
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='ada
m')

    return model

```

In [0]:

```

# https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-p
ython-keras/

activ = ['sigmoid','relu']

from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV

model = KerasClassifier(build_fn=best_hyperparameters, epochs=nb_epoch, batch_size=batc
h_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter

grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result = grid.fit(X_train, Y_train)

```

In [0]:

```

print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

```

```

Best: 0.975633 using {'activ': 'relu'}
0.974650 (0.001138) with: {'activ': 'sigmoid'}
0.975633 (0.002812) with: {'activ': 'relu'}

```

5.0 Assignment

1. Please try out models with different architectures & you can experiment with number of layers as well.
2. Include error plots.
3. Compare all your models in a tabular format using prettytable library or similar ones.

5.1 Defining model parameters

In [12]:

```
output_dim = 10
input_dim = X_train.shape[1]

batch_size = 200
nb_epoch = 30
```

In [13]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
from keras.models import Sequential
from keras.layers import Dense, Activation
from datetime import datetime
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

5.2 Two hidden layers with 784-400-80-10 architecture

5.2.1 Without Drop-out & Batch Normalization

In [16]:

```
import warnings
warnings.filterwarnings("ignore")

start = datetime.now()

model1 = Sequential()
model1.add(Dense(400, activation='relu', input_shape=(input_dim,), kernel_initializer=
'he_normal'))
model1.add(Dense(80, activation='relu', kernel_initializer= 'he_normal'))
model1.add(Dense(output_dim, activation='softmax'))

print(model1.summary())

model1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
1, validation_data=(X_test, Y_test))

print('Time taken :', datetime.now() - start)
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 400)	314000
dense_8 (Dense)	(None, 80)	32080
dense_9 (Dense)	(None, 10)	810

Total params: 346,890

Trainable params: 346,890

Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/30

60000/60000 [=====] - 7s 110us/step - loss: 0.288

5 - accuracy: 0.9172 - val_loss: 0.1520 - val_accuracy: 0.9546

Epoch 2/30

60000/60000 [=====] - 6s 94us/step - loss: 0.1071

- accuracy: 0.9677 - val_loss: 0.0921 - val_accuracy: 0.9706

Epoch 3/30

60000/60000 [=====] - 6s 99us/step - loss: 0.0675

- accuracy: 0.9795 - val_loss: 0.0887 - val_accuracy: 0.9718

Epoch 4/30

60000/60000 [=====] - 6s 98us/step - loss: 0.0487

- accuracy: 0.9850 - val_loss: 0.0709 - val_accuracy: 0.9784

Epoch 5/30

60000/60000 [=====] - 6s 97us/step - loss: 0.0346

- accuracy: 0.9891 - val_loss: 0.0723 - val_accuracy: 0.9783

Epoch 6/30

60000/60000 [=====] - 6s 100us/step - loss: 0.025

6 - accuracy: 0.9919 - val_loss: 0.0729 - val_accuracy: 0.9776

Epoch 7/30

60000/60000 [=====] - 6s 100us/step - loss: 0.018

7 - accuracy: 0.9947 - val_loss: 0.0734 - val_accuracy: 0.9785

Epoch 8/30

60000/60000 [=====] - 6s 106us/step - loss: 0.015

2 - accuracy: 0.9955 - val_loss: 0.0614 - val_accuracy: 0.9808

Epoch 9/30

60000/60000 [=====] - 6s 101us/step - loss: 0.011

3 - accuracy: 0.9969 - val_loss: 0.0668 - val_accuracy: 0.9818

Epoch 10/30

60000/60000 [=====] - 6s 105us/step - loss: 0.007

7 - accuracy: 0.9981 - val_loss: 0.0715 - val_accuracy: 0.9797

Epoch 11/30

60000/60000 [=====] - 6s 96us/step - loss: 0.0098

- accuracy: 0.9968 - val_loss: 0.0817 - val_accuracy: 0.9781

Epoch 12/30

60000/60000 [=====] - 6s 103us/step - loss: 0.014

3 - accuracy: 0.9948 - val_loss: 0.0708 - val_accuracy: 0.9803

Epoch 13/30

60000/60000 [=====] - 6s 104us/step - loss: 0.006

7 - accuracy: 0.9982 - val_loss: 0.0892 - val_accuracy: 0.9784

Epoch 14/30

60000/60000 [=====] - 6s 96us/step - loss: 0.0068

- accuracy: 0.9979 - val_loss: 0.0780 - val_accuracy: 0.9811

Epoch 15/30

60000/60000 [=====] - 6s 108us/step - loss: 0.004

8 - accuracy: 0.9984 - val_loss: 0.0870 - val_accuracy: 0.9804

```
Epoch 16/30
60000/60000 [=====] - 6s 95us/step - loss: 0.0070
- accuracy: 0.9977 - val_loss: 0.0965 - val_accuracy: 0.9780
Epoch 17/30
60000/60000 [=====] - 6s 102us/step - loss: 0.008
5 - accuracy: 0.9969 - val_loss: 0.0807 - val_accuracy: 0.9808
Epoch 18/30
60000/60000 [=====] - 6s 98us/step - loss: 0.0090
- accuracy: 0.9969 - val_loss: 0.0851 - val_accuracy: 0.9814
Epoch 19/30
60000/60000 [=====] - 6s 96us/step - loss: 0.0071
- accuracy: 0.9978 - val_loss: 0.0835 - val_accuracy: 0.9808
Epoch 20/30
60000/60000 [=====] - 6s 94us/step - loss: 0.0046
- accuracy: 0.9987 - val_loss: 0.0930 - val_accuracy: 0.9803
Epoch 21/30
60000/60000 [=====] - 6s 99us/step - loss: 0.0047
- accuracy: 0.9985 - val_loss: 0.0869 - val_accuracy: 0.9822
Epoch 22/30
60000/60000 [=====] - 6s 105us/step - loss: 0.002
5 - accuracy: 0.9994 - val_loss: 0.0898 - val_accuracy: 0.9817
Epoch 23/30
60000/60000 [=====] - 6s 107us/step - loss: 0.002
7 - accuracy: 0.9991 - val_loss: 0.1400 - val_accuracy: 0.9744
Epoch 24/30
60000/60000 [=====] - 6s 98us/step - loss: 0.0083
- accuracy: 0.9971 - val_loss: 0.1122 - val_accuracy: 0.9784
Epoch 25/30
60000/60000 [=====] - 6s 102us/step - loss: 0.006
8 - accuracy: 0.9977 - val_loss: 0.0903 - val_accuracy: 0.9824
Epoch 26/30
60000/60000 [=====] - 6s 99us/step - loss: 0.0076
- accuracy: 0.9975 - val_loss: 0.0958 - val_accuracy: 0.9822
Epoch 27/30
60000/60000 [=====] - 7s 110us/step - loss: 0.002
6 - accuracy: 0.9992 - val_loss: 0.0853 - val_accuracy: 0.9839
Epoch 28/30
60000/60000 [=====] - 6s 107us/step - loss: 5.944
6e-04 - accuracy: 0.9999 - val_loss: 0.0903 - val_accuracy: 0.9834
Epoch 29/30
60000/60000 [=====] - 6s 97us/step - loss: 1.5903
e-04 - accuracy: 1.0000 - val_loss: 0.0874 - val_accuracy: 0.9844
Epoch 30/30
60000/60000 [=====] - 6s 97us/step - loss: 4.9647
e-05 - accuracy: 1.0000 - val_loss: 0.0875 - val_accuracy: 0.9848
Time taken : 0:03:02.169499
```

In [17]:

```
score1 = model1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score1[0])
print('Test accuracy:', score1[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08754952983312075
 Test accuracy: 0.9847999811172485



In [18]:


```
w_after = model1.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5.2.2 With Drop-out & Batch Normalization

In [21]:

```
import warnings
warnings.filterwarnings("ignore")
start = datetime.now()

from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout

model2 = Sequential()

model2.add(Dense(400, activation='relu', input_shape=(input_dim,), kernel_initializer=
'he_normal'))
model2.add(BatchNormalization())
model2.add(Dropout(0.5))

model2.add(Dense(80, activation='relu', kernel_initializer= 'he_normal'))
model2.add(BatchNormalization())
model2.add(Dropout(0.5))

model2.add(Dense(output_dim, activation='softmax'))

print(model2.summary())

model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
1, validation_data=(X_test, Y_test))

print('Time taken :', datetime.now() - start)
```


Model: "sequential_6"

Layer (type)	Output Shape	Param #
dense_12 (Dense)	(None, 400)	314000
batch_normalization_2 (Batch Normalization)	(None, 400)	1600
dropout_1 (Dropout)	(None, 400)	0
dense_13 (Dense)	(None, 80)	32080
batch_normalization_3 (Batch Normalization)	(None, 80)	320
dropout_2 (Dropout)	(None, 80)	0
dense_14 (Dense)	(None, 10)	810
Total params: 348,810		
Trainable params: 347,850		
Non-trainable params: 960		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/30

60000/60000 [=====] - 12s 203us/step - loss: 0.53

73 - accuracy: 0.8386 - val_loss: 0.1991 - val_accuracy: 0.9438

Epoch 2/30

60000/60000 [=====] - 9s 153us/step - loss: 0.245

0 - accuracy: 0.9280 - val_loss: 0.1270 - val_accuracy: 0.9627

Epoch 3/30

60000/60000 [=====] - 10s 162us/step - loss: 0.19

25 - accuracy: 0.9435 - val_loss: 0.1023 - val_accuracy: 0.9683

Epoch 4/30

60000/60000 [=====] - 9s 157us/step - loss: 0.153

7 - accuracy: 0.9549 - val_loss: 0.0974 - val_accuracy: 0.9716

Epoch 5/30

60000/60000 [=====] - 10s 161us/step - loss: 0.13

72 - accuracy: 0.9600 - val_loss: 0.0774 - val_accuracy: 0.9756

Epoch 6/30

60000/60000 [=====] - 9s 158us/step - loss: 0.121

1 - accuracy: 0.9636 - val_loss: 0.0734 - val_accuracy: 0.9761

Epoch 7/30

60000/60000 [=====] - 10s 165us/step - loss: 0.11

10 - accuracy: 0.9676 - val_loss: 0.0746 - val_accuracy: 0.9763

Epoch 8/30

60000/60000 [=====] - 9s 147us/step - loss: 0.103

4 - accuracy: 0.9682 - val_loss: 0.0673 - val_accuracy: 0.9788

Epoch 9/30

60000/60000 [=====] - 9s 143us/step - loss: 0.098

2 - accuracy: 0.9704 - val_loss: 0.0661 - val_accuracy: 0.9812

Epoch 10/30

60000/60000 [=====] - 8s 142us/step - loss: 0.095

2 - accuracy: 0.9704 - val_loss: 0.0664 - val_accuracy: 0.9799

Epoch 11/30

60000/60000 [=====] - 9s 149us/step - loss: 0.087

2 - accuracy: 0.9732 - val_loss: 0.0652 - val_accuracy: 0.9793

Epoch 12/30

60000/60000 [=====] - 9s 148us/step - loss: 0.078

6 - accuracy: 0.9759 - val_loss: 0.0644 - val_accuracy: 0.9799

Epoch 13/30

```
60000/60000 [=====] - 9s 143us/step - loss: 0.078
4 - accuracy: 0.9764 - val_loss: 0.0650 - val_accuracy: 0.9798
Epoch 14/30
60000/60000 [=====] - 9s 142us/step - loss: 0.076
3 - accuracy: 0.9770 - val_loss: 0.0617 - val_accuracy: 0.9822
Epoch 15/30
60000/60000 [=====] - 9s 150us/step - loss: 0.074
0 - accuracy: 0.9773 - val_loss: 0.0658 - val_accuracy: 0.9817
Epoch 16/30
60000/60000 [=====] - 9s 145us/step - loss: 0.071
1 - accuracy: 0.9783 - val_loss: 0.0639 - val_accuracy: 0.9807
Epoch 17/30
60000/60000 [=====] - 9s 152us/step - loss: 0.066
5 - accuracy: 0.9794 - val_loss: 0.0582 - val_accuracy: 0.9828
Epoch 18/30
60000/60000 [=====] - 9s 144us/step - loss: 0.063
4 - accuracy: 0.9805 - val_loss: 0.0650 - val_accuracy: 0.9807
Epoch 19/30
60000/60000 [=====] - 9s 143us/step - loss: 0.060
7 - accuracy: 0.9804 - val_loss: 0.0592 - val_accuracy: 0.9829
Epoch 20/30
60000/60000 [=====] - 8s 137us/step - loss: 0.058
6 - accuracy: 0.9815 - val_loss: 0.0594 - val_accuracy: 0.9825
Epoch 21/30
60000/60000 [=====] - 8s 135us/step - loss: 0.057
3 - accuracy: 0.9818 - val_loss: 0.0630 - val_accuracy: 0.9826
Epoch 22/30
60000/60000 [=====] - 9s 142us/step - loss: 0.059
2 - accuracy: 0.9808 - val_loss: 0.0620 - val_accuracy: 0.9820
Epoch 23/30
60000/60000 [=====] - 9s 143us/step - loss: 0.054
4 - accuracy: 0.9827 - val_loss: 0.0598 - val_accuracy: 0.9839
Epoch 24/30
60000/60000 [=====] - 8s 141us/step - loss: 0.053
2 - accuracy: 0.9830 - val_loss: 0.0552 - val_accuracy: 0.9847
Epoch 25/30
60000/60000 [=====] - 9s 145us/step - loss: 0.051
2 - accuracy: 0.9844 - val_loss: 0.0608 - val_accuracy: 0.9832
Epoch 26/30
60000/60000 [=====] - 8s 140us/step - loss: 0.047
7 - accuracy: 0.9849 - val_loss: 0.0556 - val_accuracy: 0.9851
Epoch 27/30
60000/60000 [=====] - 8s 136us/step - loss: 0.046
6 - accuracy: 0.9855 - val_loss: 0.0564 - val_accuracy: 0.9850
Epoch 28/30
60000/60000 [=====] - 9s 143us/step - loss: 0.047
8 - accuracy: 0.9851 - val_loss: 0.0561 - val_accuracy: 0.9843
Epoch 29/30
60000/60000 [=====] - 8s 140us/step - loss: 0.050
0 - accuracy: 0.9846 - val_loss: 0.0589 - val_accuracy: 0.9834
Epoch 30/30
60000/60000 [=====] - 8s 132us/step - loss: 0.046
7 - accuracy: 0.9853 - val_loss: 0.0637 - val_accuracy: 0.9825
Time taken : 0:04:28.345052
```

In [22]:

```
score2 = model2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score2[0])
print('Test accuracy:', score2[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06374456858527555

Test accuracy: 0.9825000166893005



In [23]:


```
w_after = model2.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5.3 Three hidden layers with 784-450-200-90-10 architecture

5.3.1 Without Drop-out & Batch Normalization

In [24]:

```
import warnings
warnings.filterwarnings("ignore")
start = datetime.now()

model3 = Sequential()
model3.add(Dense(450, activation='relu', input_shape=(input_dim,), kernel_initializer=
'he_normal'))
model3.add(Dense(200, activation='relu', kernel_initializer= 'he_normal'))
model3.add(Dense(90, activation='relu', kernel_initializer= 'he_normal'))
model3.add(Dense(output_dim, activation='softmax'))

print(model3.summary())

model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
1, validation_data=(X_test, Y_test))

print('Time taken :', datetime.now() - start)
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
dense_15 (Dense)	(None, 450)	353250
dense_16 (Dense)	(None, 200)	90200
dense_17 (Dense)	(None, 90)	18090
dense_18 (Dense)	(None, 10)	910
Total params: 462,450		
Trainable params: 462,450		
Non-trainable params: 0		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/30

60000/60000 [=====] - 9s 144us/step - loss: 0.258

8 - accuracy: 0.9239 - val_loss: 0.1199 - val_accuracy: 0.9639

Epoch 2/30

60000/60000 [=====] - 7s 111us/step - loss: 0.090

7 - accuracy: 0.9724 - val_loss: 0.0888 - val_accuracy: 0.9722

Epoch 3/30

60000/60000 [=====] - 7s 121us/step - loss: 0.059

8 - accuracy: 0.9811 - val_loss: 0.0710 - val_accuracy: 0.9777

Epoch 4/30

60000/60000 [=====] - 7s 122us/step - loss: 0.040

6 - accuracy: 0.9878 - val_loss: 0.0760 - val_accuracy: 0.9776

Epoch 5/30

60000/60000 [=====] - 8s 127us/step - loss: 0.029

7 - accuracy: 0.9902 - val_loss: 0.0729 - val_accuracy: 0.9777

Epoch 6/30

60000/60000 [=====] - 7s 123us/step - loss: 0.023

0 - accuracy: 0.9926 - val_loss: 0.0777 - val_accuracy: 0.9782

Epoch 7/30

60000/60000 [=====] - 7s 122us/step - loss: 0.020

1 - accuracy: 0.9933 - val_loss: 0.0789 - val_accuracy: 0.9789

Epoch 8/30

60000/60000 [=====] - 7s 116us/step - loss: 0.015

5 - accuracy: 0.9949 - val_loss: 0.0892 - val_accuracy: 0.9766

Epoch 9/30

60000/60000 [=====] - 7s 124us/step - loss: 0.015

1 - accuracy: 0.9950 - val_loss: 0.1116 - val_accuracy: 0.9732

Epoch 10/30

60000/60000 [=====] - 7s 116us/step - loss: 0.013

6 - accuracy: 0.9953 - val_loss: 0.0938 - val_accuracy: 0.9760

Epoch 11/30

60000/60000 [=====] - 7s 123us/step - loss: 0.015

3 - accuracy: 0.9948 - val_loss: 0.0877 - val_accuracy: 0.9774

Epoch 12/30

60000/60000 [=====] - 7s 113us/step - loss: 0.012

0 - accuracy: 0.9959 - val_loss: 0.0711 - val_accuracy: 0.9825

Epoch 13/30

60000/60000 [=====] - 7s 123us/step - loss: 0.008

7 - accuracy: 0.9973 - val_loss: 0.0886 - val_accuracy: 0.9799

Epoch 14/30

60000/60000 [=====] - 7s 121us/step - loss: 0.008

2 - accuracy: 0.9969 - val_loss: 0.0933 - val_accuracy: 0.9801

Epoch 15/30

```
60000/60000 [=====] - 7s 121us/step - loss: 0.010
9 - accuracy: 0.9962 - val_loss: 0.0912 - val_accuracy: 0.9790
Epoch 16/30
60000/60000 [=====] - 7s 125us/step - loss: 0.010
1 - accuracy: 0.9964 - val_loss: 0.0963 - val_accuracy: 0.9803
Epoch 17/30
60000/60000 [=====] - 7s 117us/step - loss: 0.009
6 - accuracy: 0.9967 - val_loss: 0.0965 - val_accuracy: 0.9808
Epoch 18/30
60000/60000 [=====] - 7s 114us/step - loss: 0.008
0 - accuracy: 0.9974 - val_loss: 0.0929 - val_accuracy: 0.9803
Epoch 19/30
60000/60000 [=====] - 8s 127us/step - loss: 0.006
1 - accuracy: 0.9983 - val_loss: 0.0910 - val_accuracy: 0.9816
Epoch 20/30
60000/60000 [=====] - 7s 122us/step - loss: 0.003
9 - accuracy: 0.9989 - val_loss: 0.1026 - val_accuracy: 0.9813
Epoch 21/30
60000/60000 [=====] - 7s 123us/step - loss: 0.007
7 - accuracy: 0.9976 - val_loss: 0.1104 - val_accuracy: 0.9766
Epoch 22/30
60000/60000 [=====] - 8s 128us/step - loss: 0.010
0 - accuracy: 0.9969 - val_loss: 0.0963 - val_accuracy: 0.9808
Epoch 23/30
60000/60000 [=====] - 8s 130us/step - loss: 0.009
2 - accuracy: 0.9969 - val_loss: 0.0904 - val_accuracy: 0.9812
Epoch 24/30
60000/60000 [=====] - 8s 129us/step - loss: 0.007
3 - accuracy: 0.9978 - val_loss: 0.0994 - val_accuracy: 0.9793
Epoch 25/30
60000/60000 [=====] - 8s 125us/step - loss: 0.004
3 - accuracy: 0.9987 - val_loss: 0.0971 - val_accuracy: 0.9809
Epoch 26/30
60000/60000 [=====] - 8s 139us/step - loss: 0.003
1 - accuracy: 0.9991 - val_loss: 0.1183 - val_accuracy: 0.9784
Epoch 27/30
60000/60000 [=====] - 8s 127us/step - loss: 0.009
8 - accuracy: 0.9968 - val_loss: 0.1065 - val_accuracy: 0.9797
Epoch 28/30
60000/60000 [=====] - 7s 122us/step - loss: 0.004
5 - accuracy: 0.9983 - val_loss: 0.0967 - val_accuracy: 0.9814
Epoch 29/30
60000/60000 [=====] - 8s 128us/step - loss: 0.003
3 - accuracy: 0.9989 - val_loss: 0.1072 - val_accuracy: 0.9805
Epoch 30/30
60000/60000 [=====] - 8s 127us/step - loss: 0.009
4 - accuracy: 0.9970 - val_loss: 0.0883 - val_accuracy: 0.9821
Time taken : 0:03:44.069489
```

In [25]:


```
score3 = model3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score3[0])
print('Test accuracy:', score3[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08829341124333276

Test accuracy: 0.9821000099182129



In [26]:


```
w_after = model3.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5.3.2 With Drop-out & Batch Normalization

In [27]:

```
import warnings
warnings.filterwarnings("ignore")
start = datetime.now()

model4 = Sequential()

model4.add(Dense(450, activation='relu', input_shape=(input_dim,), kernel_initializer=
'he_normal'))
model4.add(BatchNormalization())
model4.add(Dropout(0.5))

model4.add(Dense(200, activation='relu', kernel_initializer= 'he_normal'))
model4.add(BatchNormalization())
model4.add(Dropout(0.5))

model4.add(Dense(90, activation='relu', kernel_initializer= 'he_normal'))
model4.add(BatchNormalization())
model4.add(Dropout(0.5))

model4.add(Dense(output_dim, activation='softmax'))

print(model4.summary())

model4.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model4.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
1, validation_data=(X_test, Y_test))

print('Time taken :', datetime.now() - start)
```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
dense_19 (Dense)	(None, 450)	353250
batch_normalization_4 (Batch Normalization)	(None, 450)	1800
dropout_3 (Dropout)	(None, 450)	0
dense_20 (Dense)	(None, 200)	90200
batch_normalization_5 (Batch Normalization)	(None, 200)	800
dropout_4 (Dropout)	(None, 200)	0
dense_21 (Dense)	(None, 90)	18090
batch_normalization_6 (Batch Normalization)	(None, 90)	360
dropout_5 (Dropout)	(None, 90)	0
dense_22 (Dense)	(None, 10)	910
Total params: 465,410		
Trainable params: 463,930		
Non-trainable params: 1,480		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/30

60000/60000 [=====] - 15s 254us/step - loss: 0.71

32 - accuracy: 0.7818 - val_loss: 0.2050 - val_accuracy: 0.9390

Epoch 2/30

60000/60000 [=====] - 12s 202us/step - loss: 0.28

98 - accuracy: 0.9158 - val_loss: 0.1404 - val_accuracy: 0.9568

Epoch 3/30

60000/60000 [=====] - 12s 202us/step - loss: 0.22

14 - accuracy: 0.9354 - val_loss: 0.1188 - val_accuracy: 0.9636

Epoch 4/30

60000/60000 [=====] - 12s 205us/step - loss: 0.18

19 - accuracy: 0.9466 - val_loss: 0.0996 - val_accuracy: 0.9690

Epoch 5/30

60000/60000 [=====] - 12s 206us/step - loss: 0.15

91 - accuracy: 0.9540 - val_loss: 0.0930 - val_accuracy: 0.9725

Epoch 6/30

60000/60000 [=====] - 12s 201us/step - loss: 0.14

23 - accuracy: 0.9592 - val_loss: 0.0862 - val_accuracy: 0.9727

Epoch 7/30

60000/60000 [=====] - 12s 206us/step - loss: 0.13

03 - accuracy: 0.9620 - val_loss: 0.0850 - val_accuracy: 0.9749

Epoch 8/30

60000/60000 [=====] - 12s 208us/step - loss: 0.12

19 - accuracy: 0.9645 - val_loss: 0.0767 - val_accuracy: 0.9771

Epoch 9/30

60000/60000 [=====] - 12s 199us/step - loss: 0.11

05 - accuracy: 0.9673 - val_loss: 0.0790 - val_accuracy: 0.9772

Epoch 10/30

60000/60000 [=====] - 12s 203us/step - loss: 0.10

25 - accuracy: 0.9694 - val_loss: 0.0763 - val_accuracy: 0.9785

Epoch 11/30

```
60000/60000 [=====] - 12s 197us/step - loss: 0.10
02 - accuracy: 0.9707 - val_loss: 0.0732 - val_accuracy: 0.9799
Epoch 12/30
60000/60000 [=====] - 12s 204us/step - loss: 0.09
47 - accuracy: 0.9721 - val_loss: 0.0720 - val_accuracy: 0.9796
Epoch 13/30
60000/60000 [=====] - 12s 199us/step - loss: 0.08
95 - accuracy: 0.9738 - val_loss: 0.0700 - val_accuracy: 0.9812
Epoch 14/30
60000/60000 [=====] - 12s 196us/step - loss: 0.08
59 - accuracy: 0.9749 - val_loss: 0.0661 - val_accuracy: 0.9810
Epoch 15/30
60000/60000 [=====] - 12s 198us/step - loss: 0.08
08 - accuracy: 0.9762 - val_loss: 0.0721 - val_accuracy: 0.9803
Epoch 16/30
60000/60000 [=====] - 12s 206us/step - loss: 0.07
73 - accuracy: 0.9765 - val_loss: 0.0718 - val_accuracy: 0.9807
Epoch 17/30
60000/60000 [=====] - 12s 200us/step - loss: 0.07
48 - accuracy: 0.9778 - val_loss: 0.0663 - val_accuracy: 0.9816
Epoch 18/30
60000/60000 [=====] - 12s 206us/step - loss: 0.07
05 - accuracy: 0.9788 - val_loss: 0.0629 - val_accuracy: 0.9825
Epoch 19/30
60000/60000 [=====] - 12s 201us/step - loss: 0.06
94 - accuracy: 0.9796 - val_loss: 0.0667 - val_accuracy: 0.9812
Epoch 20/30
60000/60000 [=====] - 12s 196us/step - loss: 0.07
05 - accuracy: 0.9787 - val_loss: 0.0765 - val_accuracy: 0.9799
Epoch 21/30
60000/60000 [=====] - 12s 204us/step - loss: 0.06
31 - accuracy: 0.9811 - val_loss: 0.0665 - val_accuracy: 0.9826
Epoch 22/30
60000/60000 [=====] - 13s 208us/step - loss: 0.06
53 - accuracy: 0.9805 - val_loss: 0.0760 - val_accuracy: 0.9799
Epoch 23/30
60000/60000 [=====] - 12s 201us/step - loss: 0.06
09 - accuracy: 0.9823 - val_loss: 0.0652 - val_accuracy: 0.9832
Epoch 24/30
60000/60000 [=====] - 12s 197us/step - loss: 0.06
01 - accuracy: 0.9818 - val_loss: 0.0641 - val_accuracy: 0.9833
Epoch 25/30
60000/60000 [=====] - 12s 203us/step - loss: 0.05
51 - accuracy: 0.9834 - val_loss: 0.0622 - val_accuracy: 0.9838
Epoch 26/30
60000/60000 [=====] - 12s 206us/step - loss: 0.05
61 - accuracy: 0.9827 - val_loss: 0.0651 - val_accuracy: 0.9826
Epoch 27/30
60000/60000 [=====] - 12s 200us/step - loss: 0.05
72 - accuracy: 0.9829 - val_loss: 0.0604 - val_accuracy: 0.9827
Epoch 28/30
60000/60000 [=====] - 12s 208us/step - loss: 0.05
04 - accuracy: 0.9844 - val_loss: 0.0640 - val_accuracy: 0.9834
Epoch 29/30
60000/60000 [=====] - 12s 202us/step - loss: 0.05
10 - accuracy: 0.9842 - val_loss: 0.0712 - val_accuracy: 0.9824
Epoch 30/30
60000/60000 [=====] - 12s 203us/step - loss: 0.04
82 - accuracy: 0.9849 - val_loss: 0.0648 - val_accuracy: 0.9834
Time taken : 0:06:09.783655
```

In [28]:


```
score4 = model4.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score4[0])
print('Test accuracy:', score4[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06478801985616738

Test accuracy: 0.9833999872207642



In [29]:


```
w_after = model4.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5.4 Five hidden layers with 784-450-300-200-150-90-10 architecture

5.4.1 Without Drop-out & Batch Normalization

In [30]:

```
import warnings
warnings.filterwarnings("ignore")
start = datetime.now()

model5 = Sequential()
model5.add(Dense(450, activation='relu', input_shape=(input_dim,), kernel_initializer=
'he_normal'))
model5.add(Dense(300, activation='relu', kernel_initializer= 'he_normal'))
model5.add(Dense(200, activation='relu', kernel_initializer= 'he_normal'))
model5.add(Dense(150, activation='relu', kernel_initializer= 'he_normal'))
model5.add(Dense(90, activation='relu', kernel_initializer= 'he_normal'))
model5.add(Dense(output_dim, activation='softmax'))

print(model5.summary())

model5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
1, validation_data=(X_test, Y_test))

print('Time taken :', datetime.now() - start)
```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
dense_23 (Dense)	(None, 450)	353250
dense_24 (Dense)	(None, 300)	135300
dense_25 (Dense)	(None, 200)	60200
dense_26 (Dense)	(None, 150)	30150
dense_27 (Dense)	(None, 90)	13590
dense_28 (Dense)	(None, 10)	910
Total params: 593,400		
Trainable params: 593,400		
Non-trainable params: 0		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/30

60000/60000 [=====] - 10s 173us/step - loss: 0.2548 - accuracy: 0.9237 - val_loss: 0.1071 - val_accuracy: 0.9677

Epoch 2/30

60000/60000 [=====] - 9s 149us/step - loss: 0.0917 - accuracy: 0.9717 - val_loss: 0.0830 - val_accuracy: 0.9737

Epoch 3/30

60000/60000 [=====] - 10s 160us/step - loss: 0.0610 - accuracy: 0.9811 - val_loss: 0.0746 - val_accuracy: 0.9760

Epoch 4/30

60000/60000 [=====] - 9s 154us/step - loss: 0.0453 - accuracy: 0.9856 - val_loss: 0.0750 - val_accuracy: 0.9792

Epoch 5/30

60000/60000 [=====] - 10s 158us/step - loss: 0.0355 - accuracy: 0.9887 - val_loss: 0.0727 - val_accuracy: 0.9797

Epoch 6/30

60000/60000 [=====] - 10s 161us/step - loss: 0.0275 - accuracy: 0.9910 - val_loss: 0.0745 - val_accuracy: 0.9789

Epoch 7/30

60000/60000 [=====] - 9s 156us/step - loss: 0.0231 - accuracy: 0.9924 - val_loss: 0.0933 - val_accuracy: 0.9763

Epoch 8/30

60000/60000 [=====] - 9s 157us/step - loss: 0.0228 - accuracy: 0.9927 - val_loss: 0.0961 - val_accuracy: 0.9749

Epoch 9/30

60000/60000 [=====] - 9s 158us/step - loss: 0.0220 - accuracy: 0.9920 - val_loss: 0.0920 - val_accuracy: 0.9770

Epoch 10/30

60000/60000 [=====] - 9s 156us/step - loss: 0.0188 - accuracy: 0.9940 - val_loss: 0.0741 - val_accuracy: 0.9814

Epoch 11/30

60000/60000 [=====] - 9s 155us/step - loss: 0.0163 - accuracy: 0.9948 - val_loss: 0.1015 - val_accuracy: 0.9748

Epoch 12/30

60000/60000 [=====] - 9s 150us/step - loss: 0.0143 - accuracy: 0.9956 - val_loss: 0.0974 - val_accuracy: 0.9777

Epoch 13/30

60000/60000 [=====] - 9s 153us/step - loss: 0.0169 - accuracy: 0.9947 - val_loss: 0.0889 - val_accuracy: 0.9810


```
Epoch 14/30
60000/60000 [=====] - 10s 159us/step - loss: 0.01
35 - accuracy: 0.9957 - val_loss: 0.0894 - val_accuracy: 0.9790
Epoch 15/30
60000/60000 [=====] - 9s 156us/step - loss: 0.014
0 - accuracy: 0.9955 - val_loss: 0.0675 - val_accuracy: 0.9837
Epoch 16/30
60000/60000 [=====] - 9s 150us/step - loss: 0.009
0 - accuracy: 0.9972 - val_loss: 0.0977 - val_accuracy: 0.9792
Epoch 17/30
60000/60000 [=====] - 9s 154us/step - loss: 0.010
6 - accuracy: 0.9968 - val_loss: 0.0918 - val_accuracy: 0.9818
Epoch 18/30
60000/60000 [=====] - 9s 158us/step - loss: 0.013
9 - accuracy: 0.9960 - val_loss: 0.1033 - val_accuracy: 0.9759
Epoch 19/30
60000/60000 [=====] - 9s 154us/step - loss: 0.007
9 - accuracy: 0.9975 - val_loss: 0.0980 - val_accuracy: 0.9793
Epoch 20/30
60000/60000 [=====] - 9s 151us/step - loss: 0.012
1 - accuracy: 0.9963 - val_loss: 0.0815 - val_accuracy: 0.9805
Epoch 21/30
60000/60000 [=====] - 9s 157us/step - loss: 0.009
7 - accuracy: 0.9971 - val_loss: 0.0908 - val_accuracy: 0.9824
Epoch 22/30
60000/60000 [=====] - 9s 145us/step - loss: 0.010
8 - accuracy: 0.9966 - val_loss: 0.0931 - val_accuracy: 0.9824
Epoch 23/30
60000/60000 [=====] - 9s 154us/step - loss: 0.009
5 - accuracy: 0.9972 - val_loss: 0.0792 - val_accuracy: 0.9826
Epoch 24/30
60000/60000 [=====] - 10s 159us/step - loss: 0.00
59 - accuracy: 0.9984 - val_loss: 0.0898 - val_accuracy: 0.9820
Epoch 25/30
60000/60000 [=====] - 9s 156us/step - loss: 0.007
9 - accuracy: 0.9979 - val_loss: 0.0834 - val_accuracy: 0.9827
Epoch 26/30
60000/60000 [=====] - 9s 155us/step - loss: 0.008
3 - accuracy: 0.9973 - val_loss: 0.1023 - val_accuracy: 0.9785
Epoch 27/30
60000/60000 [=====] - 9s 157us/step - loss: 0.007
4 - accuracy: 0.9979 - val_loss: 0.1035 - val_accuracy: 0.9775
Epoch 28/30
60000/60000 [=====] - 10s 166us/step - loss: 0.01
19 - accuracy: 0.9967 - val_loss: 0.0993 - val_accuracy: 0.9802
Epoch 29/30
60000/60000 [=====] - 10s 162us/step - loss: 0.00
53 - accuracy: 0.9984 - val_loss: 0.0920 - val_accuracy: 0.9821
Epoch 30/30
60000/60000 [=====] - 10s 160us/step - loss: 0.00
80 - accuracy: 0.9975 - val_loss: 0.0884 - val_accuracy: 0.9811
Time taken : 0:04:43.392793
```

In [31]:

```
score5 = model5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score5[0])
print('Test accuracy:', score5[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0884263347170416

Test accuracy: 0.9811000227928162



In [32]:


```
w_after = model5.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5.4.2 With Drop-out & Batch Normalization

In [33]:

```
import warnings
warnings.filterwarnings("ignore")
start = datetime.now()

model6 = Sequential()

model6.add(Dense(450, activation='relu', input_shape=(input_dim,), kernel_initializer=
'he_normal'))
model6.add(BatchNormalization())
model6.add(Dropout(0.5))

model6.add(Dense(300, activation='relu', kernel_initializer= 'he_normal'))
model6.add(BatchNormalization())
model6.add(Dropout(0.5))

model6.add(Dense(200, activation='relu', kernel_initializer= 'he_normal'))
model6.add(BatchNormalization())
model6.add(Dropout(0.5))

model6.add(Dense(150, activation='relu', kernel_initializer= 'he_normal'))
model6.add(BatchNormalization())
model6.add(Dropout(0.5))

model6.add(Dense(90, activation='relu', kernel_initializer= 'he_normal'))
model6.add(BatchNormalization())
model6.add(Dropout(0.5))

model6.add(Dense(output_dim, activation='softmax'))

print(model6.summary())

model6.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model6.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
1, validation_data=(X_test, Y_test))

print('Time taken :', datetime.now() - start)
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
dense_29 (Dense)	(None, 450)	353250
batch_normalization_7 (Batch Normalization)	(None, 450)	1800
dropout_6 (Dropout)	(None, 450)	0
dense_30 (Dense)	(None, 300)	135300
batch_normalization_8 (Batch Normalization)	(None, 300)	1200
dropout_7 (Dropout)	(None, 300)	0
dense_31 (Dense)	(None, 200)	60200
batch_normalization_9 (Batch Normalization)	(None, 200)	800
dropout_8 (Dropout)	(None, 200)	0
dense_32 (Dense)	(None, 150)	30150
batch_normalization_10 (Batch Normalization)	(None, 150)	600
dropout_9 (Dropout)	(None, 150)	0
dense_33 (Dense)	(None, 90)	13590
batch_normalization_11 (Batch Normalization)	(None, 90)	360
dropout_10 (Dropout)	(None, 90)	0
dense_34 (Dense)	(None, 10)	910
Total params: 598,160		
Trainable params: 595,780		
Non-trainable params: 2,380		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/30

60000/60000 [=====] - 23s 379us/step - loss: 1.33

28 - accuracy: 0.5824 - val_loss: 0.3538 - val_accuracy: 0.8937

Epoch 2/30

60000/60000 [=====] - 19s 311us/step - loss: 0.44

88 - accuracy: 0.8709 - val_loss: 0.1935 - val_accuracy: 0.9436

Epoch 3/30

60000/60000 [=====] - 18s 300us/step - loss: 0.31

56 - accuracy: 0.9134 - val_loss: 0.1593 - val_accuracy: 0.9554

Epoch 4/30

60000/60000 [=====] - 18s 306us/step - loss: 0.25

85 - accuracy: 0.9303 - val_loss: 0.1410 - val_accuracy: 0.9619

Epoch 5/30

60000/60000 [=====] - 19s 309us/step - loss: 0.21

53 - accuracy: 0.9427 - val_loss: 0.1274 - val_accuracy: 0.9667

Epoch 6/30

60000/60000 [=====] - 18s 294us/step - loss: 0.19

54 - accuracy: 0.9487 - val_loss: 0.1095 - val_accuracy: 0.9716

Epoch 7/30

```
60000/60000 [=====] - 18s 301us/step - loss: 0.17
31 - accuracy: 0.9539 - val_loss: 0.1009 - val_accuracy: 0.9739
Epoch 8/30
60000/60000 [=====] - 18s 301us/step - loss: 0.16
25 - accuracy: 0.9574 - val_loss: 0.1026 - val_accuracy: 0.9730
Epoch 9/30
60000/60000 [=====] - 18s 297us/step - loss: 0.14
82 - accuracy: 0.9612 - val_loss: 0.1006 - val_accuracy: 0.9738
Epoch 10/30
60000/60000 [=====] - 18s 303us/step - loss: 0.14
36 - accuracy: 0.9621 - val_loss: 0.0925 - val_accuracy: 0.9761
Epoch 11/30
60000/60000 [=====] - 18s 302us/step - loss: 0.13
76 - accuracy: 0.9634 - val_loss: 0.0919 - val_accuracy: 0.9764
Epoch 12/30
60000/60000 [=====] - 18s 297us/step - loss: 0.12
75 - accuracy: 0.9661 - val_loss: 0.0802 - val_accuracy: 0.9786
Epoch 13/30
60000/60000 [=====] - 17s 290us/step - loss: 0.12
25 - accuracy: 0.9668 - val_loss: 0.0847 - val_accuracy: 0.9774
Epoch 14/30
60000/60000 [=====] - 17s 289us/step - loss: 0.11
98 - accuracy: 0.9682 - val_loss: 0.0837 - val_accuracy: 0.9787
Epoch 15/30
60000/60000 [=====] - 18s 295us/step - loss: 0.11
35 - accuracy: 0.9698 - val_loss: 0.0853 - val_accuracy: 0.9788
Epoch 16/30
60000/60000 [=====] - 18s 302us/step - loss: 0.10
68 - accuracy: 0.9711 - val_loss: 0.0835 - val_accuracy: 0.9782
Epoch 17/30
60000/60000 [=====] - 18s 307us/step - loss: 0.10
15 - accuracy: 0.9728 - val_loss: 0.0765 - val_accuracy: 0.9815
Epoch 18/30
60000/60000 [=====] - 18s 298us/step - loss: 0.09
94 - accuracy: 0.9736 - val_loss: 0.0791 - val_accuracy: 0.9790
Epoch 19/30
60000/60000 [=====] - 18s 293us/step - loss: 0.09
75 - accuracy: 0.9744 - val_loss: 0.0767 - val_accuracy: 0.9818
Epoch 20/30
60000/60000 [=====] - 18s 306us/step - loss: 0.09
25 - accuracy: 0.9750 - val_loss: 0.0768 - val_accuracy: 0.9801
Epoch 21/30
60000/60000 [=====] - 18s 302us/step - loss: 0.08
89 - accuracy: 0.9768 - val_loss: 0.0702 - val_accuracy: 0.9823
Epoch 22/30
60000/60000 [=====] - 18s 302us/step - loss: 0.08
38 - accuracy: 0.9772 - val_loss: 0.0761 - val_accuracy: 0.9811
Epoch 23/30
60000/60000 [=====] - 18s 295us/step - loss: 0.08
67 - accuracy: 0.9776 - val_loss: 0.0685 - val_accuracy: 0.9835
Epoch 24/30
60000/60000 [=====] - 17s 289us/step - loss: 0.08
49 - accuracy: 0.9775 - val_loss: 0.0739 - val_accuracy: 0.9819
Epoch 25/30
60000/60000 [=====] - 18s 302us/step - loss: 0.07
57 - accuracy: 0.9789 - val_loss: 0.0746 - val_accuracy: 0.9810
Epoch 26/30
60000/60000 [=====] - 18s 300us/step - loss: 0.07
87 - accuracy: 0.9790 - val_loss: 0.0725 - val_accuracy: 0.9829
Epoch 27/30
60000/60000 [=====] - 17s 289us/step - loss: 0.07
```

```

66 - accuracy: 0.9791 - val_loss: 0.0664 - val_accuracy: 0.9841
Epoch 28/30
60000/60000 [=====] - 18s 297us/step - loss: 0.07
83 - accuracy: 0.9794 - val_loss: 0.0725 - val_accuracy: 0.9823
Epoch 29/30
60000/60000 [=====] - 18s 297us/step - loss: 0.07
15 - accuracy: 0.9807 - val_loss: 0.0672 - val_accuracy: 0.9829
Epoch 30/30
60000/60000 [=====] - 18s 292us/step - loss: 0.07
18 - accuracy: 0.9801 - val_loss: 0.0658 - val_accuracy: 0.9830
Time taken : 0:09:07.031936

```

In [34]:

```

score6 = model6.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score6[0])
print('Test accuracy:', score6[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.0658454667896498
Test accuracy: 0.9829999804496765

```



In [35]:


```
w_after = model6.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5.5 Summary

In [28]:

```
import pandas as pd

df= pd.DataFrame(columns=["Number of Layers","Activation function count","Test Loss(wit
hout drop-out & batch normalization)","Test Accuracy (without drop-out & batch normaliz
ation)","Test Loss (with drop-out & batch normalization)","Test Accuracy (with drop-out
& batch normalization)"],index=['I','II','III'])
df.loc['I']=[2,(400,80),0.087, 0.9848, 0.0637, 0.9825]
df.loc['II']=[3,(450,200,90),0.08882, 0.9821, 0.0648, 0.9833]
df.loc['III']=[5,(450,300,200,150,90),0.0884, 0.9811, 0.0658, 0.9830]
df
```

Out[28]:

	Number of Layers	Activation function count	Test Loss(without drop-out & batch normalization)	Test Accuracy (without drop-out & batch normalization)	Test Loss (with drop-out & batch normalization)	Test Accuracy (with drop-out & batch normalization)
I	2	(400, 80)	0.087	0.9848	0.0637	0.9825
II	3	(450, 200, 90)	0.08882	0.9821	0.0648	0.9833
III	5	(450, 300, 200, 150, 90)	0.0884	0.9811	0.0658	0.983

- From the above table it can be observed that the accuracy for all the models (except for the 1st model) was marginally high when drop-out & batch normalization was used
- Also test loss was significantly less in drop-out & batch normalization models